

GREENFAAS: Carbon-Aware Scheduling for Serverless Workloads

Abdeldjalil Ledmi^{a,*}, Makhlof Ledmi^a, Mohammed El Habib Souidi^a, Nabil Azizi^a, Abdelkarim Oucif^a,
Ferial Laassami^a, Hichem Haouassi^a, Toufik Messaoud Maarouk^a

^aDepartment of Computer Science, ICOSI Lab, Abbes Laghrour University of Khenchela, Khenchela, 40000, Algeria

Abstract

Data centres’ growing electricity footprint has motivated a body of research on carbon-aware scheduling, in which workloads are shifted in time or space to coincide with cleaner electricity. Much of the foundational carbon-aware scheduling literature, however, targets batch or long-running jobs, while recent FaaS-specific systems each address only part of the spatial, temporal, hardware, and autoscaling design space. The assumptions of the batch-oriented systems — minutes-to-hours runtime, delay tolerance, and clean suspend/resume — are profoundly violated by serverless or Function-as-a-Service (FaaS) workloads. FaaS functions execute in milliseconds to seconds, often under sub-second SLAs, suffer expensive cold starts when interrupted, and arrive in bursty, event-driven patterns.

We present GREENFAAS, a carbon-aware scheduling framework for serverless workloads that unifies three decisions previously treated separately: spatial routing across regions, temporal deferral within SLA bounds, and warm-pool-aware container reuse. Four prior FaaS-targeted carbon-aware schedulers cover different facets of the design space: *GreenCourier* (Chadha et al., 2023) routes single functions to low-carbon regions; *Caribou* (Gsteiger et al., 2024) shifts whole serverless workflows geographically; *EcoLife* (Jiang et al., 2024) exploits multi-generation hardware; and *CASA* (Qi et al., 2024) co-optimizes autoscaling and SLOs within a single cluster. GREENFAAS is the first to give a closed-form joint analysis of the spatial and temporal axes under cold-start economics, and to demonstrate empirically-validated robustness across topology and workload regimes.

The paper’s principal analytical contribution is a closed-form *cold-start carbon trade-off*: we prove a dichotomy in which the warm-pool placement decision is governed by a dimensionless idle-cost ratio $\beta = P_i T_w / (P_a \tau_e)$ relative to a critical threshold $(r - 1)(1 + \alpha)$, with a closed-form break-even condition and a unique critical arrival rate λ^* — the root of a monotone scalar equation — separating the regimes. The same idle-cost ratio governs the analogous temporal-shift trade-off with the deferral interval Δt in place of T_w , giving the scheduler principled decision rules along both axes.

We evaluate GREENFAAS in a discrete-event simulator, driven by a schema-faithful synthetic workload calibrated to the Azure Functions 2019 and 2021 characterizations and by real grid carbon-intensity data from Let’s-Wait-Awhile (Wiesner et al., 2021; dos-group, 2021) and from ElectricityMaps over the same January 2021 window as the Azure trace, against four carbon-aware baselines representative of recent work. Three findings stand out. *First*, GREENFAAS reduces operational carbon by 68–83% relative to FIFO across a range of topologies on fully-real time-aligned data (real Azure 2021 trace + real ElectricityMaps Jan 2021 carbon), and 64–72% on synthetic-workload setups. On fully-real time-aligned data the GREENFAAS–Spatial gap is 0.04 percentage points; across all topologies on real-carbon data the gap stays within 0.5pp. GREENFAAS’s distinct value over Spatial is topology-robustness rather than dominance in any single regime. *Second*, GREENFAAS preserves a *do-no-harm* property that extends beyond zero-opportunity workloads to a robust neighbourhood of low-opportunity ones: across shiftable fractions $\{0, 1\%, 2\%\}$ and in single-region scenarios on both synthetic and fully-real data, GREENFAAS matches FIFO byte-for-byte; a heuristic ablation lacking the cold-start trade-off correction loses 5–102% to FIFO in the same scenarios. *Third*, the canonical Wait-Awhile temporal scheduler and a FaaS-adapted port of the provably-optimal one-way trading algorithm of Lechowicz et al. (2023) both *fail* on FaaS workloads, losing 3–306% to FIFO depending on the regime — catastrophically (up to –306%) at batch-default deferral horizons on synthetic workloads, and mildly (3–5%) on fully-real data where they degenerate toward a FIFO no-op. In neither regime do they reduce

carbon: a clean failure mode for batch carbon-aware techniques on FaaS, including those with provable competitive guarantees, when warm-pool idle-energy coupling is not accounted for. Head-to-head against the closest competing serverless system, *Caribou* (Gsteiger et al., 2024), well-tuned Caribou (hourly re-deployment with perfect carbon forecasting) and GREENFAAS achieve statistically indistinguishable carbon savings (within 0.1pp) on real Azure 2021 workload; the difference is operational rather than empirical, with GREENFAAS avoiding the container migration and pub/sub orchestration that Caribou requires. The simulator, scheduler, trace loaders, Caribou port, and real Azure Functions and ElectricityMaps integrations are released as open-source artifacts under the Apache-2.0 license upon publication.

Keywords: carbon-aware scheduling, serverless, Function-as-a-Service, cold start, sustainable computing

1. Introduction

Data centres consume an increasing share of global electricity, and hyperscale operators have made public commitments to carbon-neutral or net-zero operation within the decade. These pledges have motivated substantial research on *carbon-aware computing*: techniques that exploit spatial and temporal variation in grid carbon intensity to shift computation toward cleaner energy. The dominant strategy in the literature is *temporal shifting* of delay-tolerant work to periods of low carbon intensity (Wiesner et al., 2021); complementary *spatial shifting* routes work toward regions with cleaner grids (Sukprasert et al., 2024; Radovanović et al., 2023). Both have delivered 20–45% operational-carbon reductions for batch processing, machine learning training, and similar long-running workloads (Hanafy et al., 2023; Sukprasert et al., 2024).

A critical assumption underlies essentially this entire body of work: the target workload is *batch* in character. The canonical scheduled job is minutes to hours long, delay-tolerant on the order of hours, amenable to clean suspend-and-resume, and accompanied by enough up-front information that the scheduler can reason globally over a set of pending jobs. These assumptions are violated profoundly by *serverless* or *Function-as-a-Service* (FaaS) workloads, which now power a substantial fraction of cloud-native applications. FaaS invocations are typically measured in tens of milliseconds to a few seconds; they are often user-facing with sub-second latency objectives; they cannot be suspended and resumed without paying the full cost of a *cold start*; their arrivals are bursty and event-driven; and per-invocation behaviour is largely unpredictable until execution.

This mismatch is consequential. Recent FaaS-targeted carbon-aware schedulers — *GreenCourier* (Chadha et al., 2023), *Caribou* (Gsteiger et al., 2024), *EcoLife* (Jiang et al., 2024), and *CASA* (Qi et al., 2024) — have each begun to address the gap from a distinct angle: spatial routing of single functions (GreenCourier), workflow-level geo-shifting (Caribou), hardware-generation choice (EcoLife), and SLO/carbon-aware autoscaling (CASA). None of them, however, formalizes the joint spatial/temporal/warm-pool decision problem that FaaS schedulers actually face, and the fundamental cold-start versus idle-energy trade-off has not previously been characterized in closed form.

We argue that carbon-aware scheduling for FaaS is both feasible and distinct enough from its batch counterpart to warrant first-class treatment. Three observations motivate our approach. *First*, FaaS workloads are heterogeneous in latency tolerance: a customer-facing API request demands sub-second response, but many event-driven functions tolerate seconds to minutes of deferral without observable user impact. *Second*, FaaS platforms already perform fine-grained placement and routing decisions for load-balancing and cold-start mitigation; carbon-aware policies can be layered onto this existing machinery. *Third*, warm-pool management — how many idle containers to keep alive in each region — is itself a carbon decision, because

*Corresponding author.

Email addresses: `abdeldjalil.ledmi@univ-khenchela.dz` (Abdeldjalil Ledmi), `ledmi.makhlouf@univ-khenchela.dz` (Makhlouf Ledmi), `souidi.mohammed@univ-khenchela.dz` (Mohammed El Habib Souidi), `azizi.nabil@univ-khenchela.dz` (Nabil Azizi), `oucif.abdelkarim@univ-khenchela.dz` (Abdelkarim Oucif), `ferial.laassami@univ-khenchela.dz` (Ferial Laassami), `haouassi.hichem@univ-khenchela.dz` (Hichem Haouassi), `maarouk.toufik@univ-khenchela.dz` (Toufik Messaoud Maarouk)

idle containers consume non-trivial power and the choice of where to keep them warm interacts directly with the carbon intensity of the host grid.

In this work, we present **GREENFAAS**, a carbon-aware scheduling framework for serverless workloads that unifies spatial routing, temporal deferral, and warm-pool-aware container reuse under a single principled policy. **GREENFAAS** combines temporal deferral for asynchronous and event-driven invocations, spatial routing for latency-tolerant classes, and carbon-aware warm-pool management, all under an SLA-tiered policy distinguishing interactive, deferrable, and background functions. A key technical contribution is the explicit modelling of the *cold-start carbon trade-off*: keeping containers warm in a low-carbon region can, in some regimes, cost more carbon (through idle power) than cold-starting in a higher-carbon region. We derive analytical break-even points for this trade-off and incorporate them directly into the scheduler.

We evaluate **GREENFAAS** in **GreenFaaS-sim**, a discrete-event simulator whose design follows the abstractions of **faas-sim** and **vessim** (Edgerun, 2022; dos-group, 2023), driven by the real Azure Functions 2021 per-invocation trace (Zhang et al., 2021; Microsoft, 2021) and by real grid carbon traces from Let’s-Wait-Awhile (Wiesner et al., 2021; dos-group, 2021) and from ElectricityMaps over the same January 2021 window as the Azure trace. The trace loaders also accept a schema-faithful synthetic workload as a fallback when the real trace is not locally available (§8.1). We compare against four carbon-aware baselines drawn from the recent literature, evaluating along five axes: workload intensity, number of available regions, SLA class mix, carbon-intensity variability, and forecast accuracy.

Our key findings are threefold. *First*, **GREENFAAS** reduces operational carbon by 68–83% across topologies on fully-real time-aligned data (Azure 2021 trace + ElectricityMaps Jan 2021 carbon), and by 64–72% on synthetic-workload setups; on fully-real data the **GREENFAAS**–Spatial gap is just 0.04 percentage points. *Second*, **GREENFAAS** preserves a *do-no-harm* property: in single-region and flat-carbon regimes where shifting offers no real opportunity, **GREENFAAS** matches FIFO byte-for-byte on both synthetic and fully-real data, while a heuristic ablation (lacking the cold-start trade-off correction) loses 5–102% to FIFO in the same scenarios. *Third*, the canonical Wait-Awhile temporal scheduler and a FaaS-adapted port of the Lechowicz et al. (2023) provably-optimal one-way trading algorithm both *fail* on FaaS workloads (3–306% over FIFO depending on the regime), an empirically sharp demonstration that batch-oriented carbon-aware techniques — including those with provable competitive guarantees — do not transfer to FaaS without addressing warm-pool idle-energy coupling. This robustness across regimes — rather than dominance in any single regime — is the central practical claim of the paper.

Contributions. The paper makes the following contributions:

1. A formal model of carbon-aware FaaS scheduling as an online optimization problem with capacity, SLA, and warm-pool constraints, and an honest discussion of why classical competitive-ratio bounds do not apply (§4).
2. A closed-form *cold-start carbon trade-off lemma* characterizing warm-pool versus cold-start carbon as a function of arrival rate, runtime and cold-start durations, active and idle power, warm-pool TTL, and the carbon-intensity ratio between regions (§5). The same dimensionless idle-cost ratio governs the analogous temporal-shift decision via an idle-energy correction (§5.4).
3. The **GREENFAAS** scheduler: a hybrid online algorithm combining temporal deferral, spatial routing, and lemma-driven warm-pool management under an SLA-tiered policy, with $O(R \cdot N)$ per-invocation cost (§6).
4. An open-source simulator with trace loaders that accept the published Azure Functions 2019, Azure Functions 2021, and LWA carbon-intensity schemas without modification (§7).
5. An empirical evaluation against four carbon-aware baselines and a five-axis sensitivity sweep establishing the do-no-harm property across topology, SLA-class mix, forecast accuracy, carbon variability, and workload intensity (§8).

The remainder of this paper is organized as follows. §2 surveys related work; §3 motivates the problem; §4 formalizes it; §5 derives the trade-off lemma; §6 presents the algorithm; §7 describes the simulator; §8 reports results; §11 discusses limitations; §12 concludes.

2. Related Work

Carbon-aware scheduling has matured rapidly over the past four years. We organise prior work into four threads: (i) batch and long-running workloads, where the dominant approaches assume delay-tolerance and suspend/resume; (ii) theoretical foundations of online carbon-aware scheduling; (iii) the small but growing set of carbon-aware schedulers explicitly targeting serverless or FaaS workloads, with which GREENFAAS most directly engages; and (iv) FaaS systems work on cold-starts, warm pools, and characterization, on which our cost model depends.

2.1. Carbon-Aware Scheduling for Batch Workloads

Wiesner et al.’s *Let’s Wait Awhile* (Wiesner et al., 2021) introduced threshold-based deferral for batch workloads and released the hourly carbon-intensity dataset (DE, GB, FR, US-CAISO, 2020) that has become the de-facto reference; we use this dataset directly in our evaluation (§8.1). *CarbonScaler* (Hanafy et al., 2023) scales the parallelism of ML training jobs in response to carbon intensity, exploiting the elasticity of data-parallel workloads. Souza et al.’s *GAIA* scheduler (Sukprasert et al., 2024) addresses the three-way trade-off between carbon, performance, and cost for batch jobs, and quantifies the practical limits of carbon-aware spatiotemporal shifting; the same group’s *Ecovisor* (Souza et al., 2023) provides a virtual energy system abstraction. Google’s Carbon-Intelligent Compute (Radovanović et al., 2023) deploys a Variable Capacity Curve approach at production scale. *LACS* (Bostandoost et al., 2024) is a learning-augmented resource scaler addressing uncertain demand. *Adapting Datacenter Capacity* (Lin and Chien, 2023) explored provisioning to track grid carbon.

These works share two structural assumptions that GREENFAAS relaxes. *First*, the schedulable units are *jobs* with cost separable per-job, allowing the schedulers to reason about each job’s deferral in isolation. FaaS invocations are not separable in this sense because their cost is coupled through warm-pool state, which we formalize in §4. *Second*, the schedulers operate on a timescale (minutes to hours) where cold-start penalties are negligible. We show empirically in §8 that naively transplanting these techniques to FaaS workloads — as the Wait-Awhile baseline does — *more than doubles* operational carbon relative to a carbon-unaware FIFO baseline (163–236% above FIFO across our synthetic and real-carbon experiments). The “batch schedulers don’t transfer” claim is therefore not only intuitively obvious from the assumption mismatch but also empirically severe.

2.2. Theoretical Foundations

A recent line of work has put carbon-aware online algorithms on firmer footing. Lechowicz et al.’s *Online Pause and Resume Problem* (Lechowicz et al., 2023) studies a one-dimensional model in which a workload can be paused during high-carbon periods and resumed during low-carbon periods, paying a fixed switching cost. They derive double-threshold algorithms with provably optimal competitive ratios. These theoretical results apply to a workload model that is in several respects strictly simpler than ours: a single workload (no multi-tenant mix), one machine (no spatial routing), and a single switching cost (no idle-energy accumulation). The FaaS setting we study introduces three structural features — per-invocation cold-start versus warm-pool trade-offs, spatial routing across regions, and tiered SLA constraints — that take us outside the class of problems for which competitive-ratio bounds are currently known. We are explicit in §4.2 that GREENFAAS is an empirical scheduler with a *local* analytical guarantee (Lemma 5.1’s break-even characterization) rather than a competitively-optimal one. Extending (Lechowicz et al., 2023) to handle warm-pool idle energy and SLA tiering is an open direction we identify in §11.

2.3. Carbon-Aware FaaS and Serverless

Five recent systems target serverless workloads specifically; each covers a different facet of the design space.

GreenCourier (Chadha et al., 2023).. The closest single-function spatial-routing precedent: a Kubernetes scheduler plugin that routes serverless functions across geographically distributed Knative clusters based on real-time grid carbon intensity. It reports 9–18% carbon reduction per invocation with negligible latency overhead. GreenCourier and GREENFAAS share the high-level goal but differ in three significant respects. *First*, GreenCourier is *spatial-only* — it routes functions to the cleanest available region, with no temporal deferral and no warm-pool reasoning. Our headline contribution beyond GreenCourier is therefore the joint spatial/temporal/warm-pool policy and the underlying trade-off analysis. *Second*, GreenCourier has no analytical model of the warm-pool versus cold-start carbon trade-off; our Lemma 5.1 gives a closed-form characterization that becomes essential as soon as the scheduler reasons about *whether* to maintain a warm pool. *Third*, GreenCourier’s reported savings (9–18%) sit at the lower end of what our headline results suggest is achievable; we attribute this to the spatial-only scope rather than to algorithmic deficiencies in GreenCourier itself.

Caribou (Gsteiger et al., 2024).. Caribou is the closest recent system in scope at the workflow level. It offloads serverless *workflows* (multi-function DAGs) across geo-distributed regions for sustainability, automatically (re-)deploying functions and redirecting traffic to lower-carbon endpoints. Caribou requires no changes to application logic and preserves performance and cost metrics. Its scope is workflow-level spatial (no temporal deferral, no per-invocation warm-pool reasoning), with decisions made at deployment time and refreshed by periodic redeployment rather than per invocation.

We compare head-to-head in §8.7 by porting Caribou’s deployment-solver logic to our per-invocation workload as a single-node DAG case, with perfect carbon forecasting and constant token-bucket grant (both favourable to Caribou). The finding: Caribou(1h) and GREENFAAS achieve carbon savings that differ by less than 0.1 percentage points on our 24h FaaS slices; the two systems are operationally distinct but empirically equivalent on per-invocation workloads with hour-scale carbon variation. The systems are complementary along the granularity axis: Caribou makes coarse-grained, deployment-time placement decisions for whole workflows; GREENFAAS makes fine-grained, per-invocation decisions including warm-pool reasoning. A natural composition would use Caribou for workflow placement and GREENFAAS for per-invocation scheduling within each chosen region; we have not implemented this composition.

EcoLife (Jiang et al., 2024).. EcoLife targets the hardware-generation axis: it co-optimizes *operational* and *embodied* carbon by exploiting multi-generation hardware (newer servers are faster but have higher embodied carbon per unit time; older servers are slower but already amortised). It uses a Particle Swarm Optimisation-based scheduler that decides, per invocation, which hardware generation to execute on. EcoLife explicitly reasons about keep-alive (warm-pool) versus cold-start trade-offs — which makes it the closest precedent for our Lemma 5.1 — but it does so through a metaheuristic optimizer without a closed-form characterization. The two contributions are largely orthogonal: EcoLife answers “on what hardware generation should this function execute?”; GREENFAAS answers “in what region and at what time should this function execute, and is a warm pool worth maintaining there?”. A scheduler combining the two axes is a natural follow-up (§11).

CASA (Qi et al., 2024).. CASA is the closest recent single-cluster, autoscaling-focused competitor. It is an SLO- and carbon-aware framework that schedules and autoscales containers in a serverless cluster, explicitly trading off the SLO cost of cold-starts against the carbon cost of pre-warming and keep-alive. CASA reports up to a 2.6× reduction in operational carbon while reducing SLO violations by up to 1.4× over a state-of-the-art baseline. The follow-on work (Qi et al., 2024) extends the same multi-objective framework to co-optimize wastewater alongside carbon and SLOs. CASA and GREENFAAS address overlapping concerns through different mechanisms. CASA’s contribution is a carbon-aware autoscaling and container-placement policy in a single cluster, evaluated by simulation against a state-of-the-art autoscaler. GREENFAAS’s contribution is a closed-form trade-off characterization (Lemma 5.1) that yields a parameter-free per-invocation decision rule, combined with a multi-region routing policy and a formal do-no-harm property (§8.8, §8.9). Where CASA establishes that carbon-aware autoscaling is worthwhile, GREENFAAS establishes the analytical conditions under which warm-pool maintenance is itself carbon-positive — a question CASA’s heuristics do not pose

explicitly. The two approaches could be composed at different granularities: CASA at the container-pool autoscaling layer, GREENFAAS at the per-invocation routing-and-warm-pool-gate layer.

GreenWhisk (Serenari et al., 2024).. GreenWhisk is a system-level redesign of Apache OpenWhisk to expose energy and carbon state to load-balancing algorithms, supporting both grid-connected and grid-isolated modes. Its contribution is infrastructural: an energy-interface API and a carbon-aware load-balancer integrated into a production-grade FaaS platform. The scheduling algorithm itself is intentionally minimal in scope; the value is the open-platform abstraction. GREENFAAS’s algorithmic and analytical contributions are compatible with this abstraction: a production GreenWhisk deployment could load Algorithm 1 as its load-balancing policy without changes to either side. GreenWhisk’s two-mode design (grid intermittency) is a deployment dimension our analysis does not address and that an integration would have to handle separately.

CASPER (Souza et al., 2024).. Adjacent rather than directly competing: targets distributed web services (MediaWiki-class applications) under latency SLOs, using a mixed-integer program to jointly provision servers and load-balance requests across cloud regions. It addresses spatial routing and provisioning with SLO constraints but treats requests at the application level (request rates, not per-invocation arrivals) and does not engage with cold-start economics. We position GREENFAAS as the per-invocation, FaaS-native analogue of CASPER, with the cold-start trade-off model as a key technical distinction.

2.4. Serverless Cold-Starts and Characterisation

Our cost model in §4 and the SLA-tiered policy in §6.3 depend on a substantial body of FaaS systems work. *Serverless in the Wild* (Shahrad et al., 2020) is the canonical characterization of production Azure Functions workloads; it established that invocations are dominated by a small number of high-frequency “hot” functions, that durations follow heavy-tailed distributions, and that arrival processes are bursty and event-driven. Our workload generator and the function popularity model (Zipf with $s = 1.2$) follow this characterization directly. The follow-up *Azure Functions Invocation Trace 2021* (Zhang et al., 2021; Microsoft, 2021) gives per-invocation records over two weeks; we use its schema as the basis for our trace loader (§8.1).

The cold-start problem itself has a rich literature. *Catalyzer* (Du et al., 2020) and *Firecracker* (Agache et al., 2020) attack cold-starts at the system level via micro-VM and snapshot techniques. *FaaS-Cache* (Fuerst and Sharma, 2021) frames warm-pool management as a caching problem and derives keep-alive policies from caching theory. We borrow the warm-pool / keep-alive abstraction directly, but add the carbon dimension: idle energy at intensity C_r in region r is itself a cost that the scheduler must trade off against cold-start energy at potentially different C . This is the substance of our Lemma 5.1.

A parallel measurement-focused line of work makes it possible to attribute per-invocation carbon footprints to FaaS workloads with high accuracy. *Accountable Carbon Footprints* (Sharma and Fuerst, 2024) develops a Shapley-value-based disaggregation methodology that fairly attributes shared hardware and software emissions to individual function invocations, achieving >99% accuracy on a wide range of workloads. This work is methodologically orthogonal — it measures, we schedule — but a production-grade carbon-aware scheduler should integrate with such a measurement substrate to close the loop from decisions to observed outcomes.

Summary of positioning. To our knowledge, no prior work has formalized the cold-start carbon trade-off in closed form, demonstrated the temporal-axis idle-energy correction (§5.4), or established a do-no-harm property across topology, SLA-mix, forecast-accuracy, and carbon-variability axes. These are the empirical and analytical novelties on which the rest of the paper builds.

3. Motivation

We motivate carbon-aware FaaS scheduling with a joint characterization of the two data sources that drive our design: a representative FaaS workload, and the grid carbon-intensity signal that determines the carbon footprint of executing that workload.




Figure 1: FaaS arrivals are decoupled from grid carbon intensity. The orange line (right axis) shows the per-minute invocation rate of the `webhook_handler` function; the green line (left axis) shows grid carbon intensity in Germany. The two signals are phase-shifted by 8–12 hours, creating temporal shifting opportunity for deferrable invocations. Shaded bands mark the lowest- (green) and highest-carbon (red) quartile windows.

3.1. Joint Characterisation

Figure 1 overlays a 48-hour grid carbon-intensity trace for Germany — a synthetic trace whose annual mean is calibrated to the ENTSO-E 2020 figure later validated against real LWA data in §8.3 — with the per-minute invocation rate of a representative `webhook_handler` function drawn from a workload generator calibrated to the Shahrad et al. (2020) Azure characterization. Shaded bands mark the lowest-quartile (green) and highest-quartile (red) carbon windows over the 48-hour horizon. Three observations emerge.

1. *FaaS arrival rates are decoupled from grid carbon intensity.* The webhook function’s rate follows a business-hours diurnal pattern, peaking in early afternoon and troughing overnight. Grid carbon intensity in Germany follows the inverse pattern: solar and wind generation peak around midday, depressing the marginal intensity, while night-time demand is served by base-load fossil generation. The two signals are phase-shifted by roughly 8–12 hours; on this representative day, 68% of webhook invocations land *outside* the cleanest carbon quartile and 18% land directly in the highest-carbon quartile. The decoupling is the structural reason carbon-aware scheduling can help.

2. *The decoupling is robust across function classes.* Across the 60-function catalog generated for our evaluation, the population-weighted invocation rate correlates only weakly with the local carbon intensity (Pearson r in the range -0.2 to $+0.1$ across regions), confirming that arrival patterns track user activity rather than grid economics — as expected, since FaaS workloads are driven by external events (user requests, message-queue depths, scheduled triggers) that have no relationship to the regional generation mix.

3. *The amplitude of the carbon swing is comparable to the amplitude of the arrival swing.* Carbon intensity in Germany swings from ~ 200 to ~ 500 gCO₂eq/kWh over the 24-hour cycle — a ratio of $2.5\times$. The webhook invocation rate swings from ~ 15 to ~ 80 invocations per minute — a ratio of $5\times$. The two signals have similar dynamic range, which means a scheduler that can shift even a modest fraction of invocations across the diurnal cycle has potential to capture meaningful savings. We quantify this potential rigorously in §8.

3.2. Implications for Scheduler Design

The characterization directly motivates three design choices in §6. *First*, the SLA-tiered policy: arrivals distributed across the carbon cycle are only useful to a scheduler that distinguishes between invocations

that must execute immediately (interactive) and those that can wait. *Second*, the use of *both* spatial and temporal axes: a purely temporal scheduler can only exploit the diurnal swing within one region; a purely spatial scheduler ignores the substantial within-region opportunity Figure 1 reveals. *Third*, the explicit carbon-aware warm-pool reasoning of §5: at the timescale of FaaS invocations (seconds), the idle-energy cost of keeping containers warm is a first-class concern.

The motivation, in short, is not “carbon-aware scheduling helps FaaS” in some generic sense — that has been shown by the four FaaS-targeted systems surveyed in §2.3 (Chadha et al., 2023; Gsteiger et al., 2024; Jiang et al., 2024; Qi et al., 2024). The motivation is that the *specific structure* of FaaS workloads (short, bursty, latency-tiered) combined with the *specific structure* of grid carbon traces (diurnal, decoupled from user activity, multi-region heterogeneous) creates a joint optimization problem that no prior scheduler has addressed in its full generality. The rest of the paper formalizes and solves that problem.

4. Problem Formulation

This section formalizes the carbon-aware FaaS scheduling problem. §4.1 sets up the entities, decision variables, and constraints. §4.2 frames the optimization as an online problem and discusses why it admits neither a tractable offline optimum nor a competitive-ratio bound, motivating the empirical evaluation strategy in §8.

4.1. Problem Statement

Regions.. A set $\mathcal{R} = \{r_1, \dots, r_R\}$ of geographic regions. Each region r has capacity K_r , a time-varying grid carbon intensity $C_r : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{> 0}$ in gCO₂eq/kWh, a Power Usage Effectiveness $\text{PUE}_r \geq 1$, and network round-trip latencies $\rho(r, r') \geq 0$ to every other region.

Functions.. A set $\mathcal{F}$ of deployed functions, each $f \in \mathcal{F}$ characterized by expected execution time $\tau_e(f)$, cold-start duration $\tau_c(f)$, active power $P_a(f)$, idle power $P_i(f)$, latency class $\ell(f) \in \{\text{INTERACTIVE}, \text{DEFERRABLE}, \text{BACKGROUND}\}$ and SLA deadline $D(f)$. The latency class determines the scheduler’s degree of freedom: INTERACTIVE invocations must execute immediately in their home region; DEFERRABLE invocations may be deferred up to a class-dependent horizon and routed within a tight RTT budget; BACKGROUND invocations may use their full deadline and any region.

Invocations.. Function invocations arrive as a stream $\mathcal{I} = \langle i_1, i_2, \dots \rangle$ ordered by arrival time. Each i specifies a function $f(i)$, an arrival time $a(i)$, a home region $h(i)$, and a realised runtime $\tau(i)$ drawn from $f(i)$ ’s execution-time distribution. Crucially, $\tau(i)$ is *not* known to the scheduler at decision time; only the expected runtime $\tau_e(f(i))$ is. The same applies to future arrivals: at time $a(i)$ the scheduler sees i but knows nothing about i_k for k with $a(i_k) > a(i)$ beyond what an arrival-rate estimator can infer.

Warm pools.. For each $(r, f) \in \mathcal{R} \times \mathcal{F}$, the platform maintains a non-negative integer warm-container count $W_{r,f}(t)$ representing containers that have previously executed f in r and remain alive awaiting reuse. A warm container is evicted after a TTL of T_w seconds of inactivity.

Decisions.. For each invocation i , the scheduler produces $\pi(i) = (r(i), s(i))$: the execution region $r(i)$ and the scheduled start time $s(i)$. The secondary outcome $w(i) = \mathbf{1}[W_{r(i),f(i)}(s(i)) > 0]$ — whether i finds a warm container at $r(i)$ at time $s(i)$ — is observed rather than chosen; if $w(i) = 0$, the invocation pays $\tau_c(f(i))$ in start-up latency and energy.

Per-invocation cost.. The carbon cost of i under $\pi(i) = (r, s)$ is

$$\text{Carbon}(i) = \frac{\text{PUE}_r}{3.6 \times 10^6} C_r(s) \cdot P_a(f(i)) [\tau(i) + (1 - w(i))\tau_c(f(i))].$$

End-to-end latency is

$$\text{Lat}(i) = (s - a(i)) + \rho(h(i), r) + (1 - w(i))\tau_c(f(i)) + \tau(i).$$

System cost.. Aggregate carbon over horizon $[0, T]$ sums per-invocation execution carbon and warm-pool idle carbon. The exact idle carbon integrates the time-varying intensity over each idle interval:

$$\begin{aligned} \text{TotalCarbon}([0, T]) = & \sum_{i: a(i) \in [0, T]} \text{Carbon}(i) \\ & + \sum_{r, f} \frac{\text{PUE}_r}{3.6 \times 10^6} P_i(f) \int_0^T W_{r, f}(t) C_r(t) dt, \end{aligned} \quad (1)$$

where $W_{r, f}(t)$ is the number of warm-but-idle containers for function f in region r at time t , and $C_r(t)$ is the instantaneous carbon intensity. For analytical tractability in Lemma 5.1 and in the SCORESLOTS heuristic, we approximate $C_r(t)$ over an idle interval by its mean $\bar{C}_r$. This approximation is exact when warm-pool idle time is distributed uniformly over the horizon (the typical case for steady FaaS load, where the warm pool is continuously cycled); deviation arises only if idle time correlates with carbon intensity. Because GreenFaaS’s spatial routing concentrates load in low-carbon regions during high-carbon periods, any residual correlation is conservative (it would slightly *over*-estimate idle carbon in the regions GreenFaaS avoids). The simulator (§7) reports results under this mean-intensity approximation; the scheduler *ranking* is what the paper claims, and the power-model robustness check (§11) confirms the ranking is insensitive to the idle-energy magnitude.

Constraints.. A schedule π must respect: (1) *Capacity*: at every t , the count of in-flight invocations in region r is at most K_r . (2) *SLA-feasibility in expectation*: $(s(i) - a(i)) + \rho(h(i), r(i)) + (1 - \hat{w}(i))\tau_c(f(i)) + \tau_e(f(i)) \leq D(f(i))$, where $\hat{w}(i)$ is the scheduler’s warm-pool prediction. (3) *Class-dependent shifting limits*: INTERACTIVE $\Rightarrow s(i) = a(i)$ and $r(i) = h(i)$; DEFERRABLE $\Rightarrow s(i) \leq a(i) + \Delta^D$ and $\rho(h(i), r(i)) \leq \rho^D$; BACKGROUND $\Rightarrow$ analogous with class-specific bounds Δ^B, ρ^B . (4) *Causality*: $s(i) \geq a(i)$.

Objective.. Minimise total carbon over $[0, T]$:

$$\pi^* = \arg \min_{\pi} \text{TotalCarbon}([0, T]; \pi) \quad \text{s.t. (1)–(4)}. \quad (2)$$

4.2. Online Optimisation Framing

Two structural features distinguish this problem from prior carbon-aware scheduling work.

Online, non-clairvoyant.. Decisions $\pi(i)$ must be committed at time $a(i)$, before any subsequent invocation is observed. The scheduler has access to past arrivals, the current state $\{W_{r, f}(t), K_r - \text{in-flight}_r(t)\}$, and a forecast $\hat{C}_r$ of future carbon intensity over a finite horizon (empirically up to 24h with degrading accuracy; see §8.12). It does *not* see future arrivals. We further restrict to *non-clairvoyant* schedulers: the realised runtime $\tau(i)$ is not revealed until execution completes; the scheduler uses $\tau_e(f(i))$ as its working estimate. Prior carbon-aware schedulers for batch workloads sometimes assume per-job runtime is known up front (Hanafy et al., 2023); we cannot.

Non-additive cost structure.. The system cost (Eq. 1) is *not* separable into per-invocation contributions because of the warm-pool idle-energy term. A scheduler’s decision $\pi(i)$ affects not only $\text{Carbon}(i)$ but also $W_{r, f}(t)$ for all $t \geq s(i)$, which in turn governs whether subsequent invocations of f in r pay a cold start and how much idle energy accrues. Decisions are thus *coupled across invocations* through the warm-pool state. This is the structural reason why batch-oriented online algorithms — which typically assume separable per-job costs — do not directly apply.

Why no competitive-ratio bound. Even with full knowledge of future arrivals, the offline problem contains generalized-assignment-like structure: invocations must be mapped to (region, time-slot) bins under per-region capacity constraints K_r and per-function deadline constraints $D(f)$, with a bin-dependent carbon cost. The classical generalized assignment problem is **NP**-hard, and the non-separable idle-energy term makes our cost function supermodular rather than linear in the assignment, so we expect the offline optimum to be at least as hard; we do not give a formal reduction, since the offline problem is not the focus of our contribution. We do not pursue an approximation-ratio analysis for two reasons. *First*, the cost function combines additive execution carbon with a non-separable idle-energy integral, which puts the problem outside the standard frameworks for online competitive analysis (makespan, machine scheduling, k -server). *Second*, an arbitrarily small adversarial perturbation of the carbon-intensity functions C_r can make any online scheduler arbitrarily worse than the offline optimum, so any finite competitive ratio would have to assume a specific stochastic model of C_r — and the empirical literature is clear that real grid carbon is neither stationary nor i.i.d. across regions. We therefore adopt an *empirical* evaluation strategy (§8); the analytical contribution of the paper is the *local* trade-off characterization of §5, which gives the scheduler a principled decision rule without claiming global optimality.

Decomposition. GREENFAAS decomposes the per-invocation problem into two sub-problems tractable individually. *Region gating* (long-run, structural): given the function’s arrival rate $\hat{\lambda}(f, h)$ and current carbon intensities, which subset of regions could ever be a carbon-cheaper warm-pool host than the home region? Lemma 5.1 (§5) answers this in closed form. *Per-slot scoring* (local, this-invocation): given the admitted regions and a deferral horizon, which (region, start-time) pair minimises expected carbon for this single invocation, charged appropriately for idle-energy effects (§5.4)? §6 makes the algorithm and its complexity precise.

5. The Cold-Start Carbon Trade-off

In this section we present the paper’s central analytical contribution: a closed-form characterization of when it is carbon-cheaper to keep a warm pool in a low-carbon region than to cold-start in a high-carbon region. The result is non-obvious because warm pools consume non-trivial idle energy; at low arrival rates that idle energy — accumulated in the low-carbon region — can exceed the cold-start penalty paid in the high-carbon region. The empirical existence of this trade-off has been noted concurrently (Jiang et al., 2024), but it has not previously been characterized in closed form, which is essential for a scheduler that needs a parameter-free rule generalizing across function profiles and region pairs.

5.1. Setup and Per-Invocation Energy

Consider a single function with execution time τ_e , cold-start duration τ_c , active power P_a (during execution and cold start), and warm-idle power P_i . Warm containers are evicted after T_w seconds. Invocations arrive as a homogeneous Poisson process with rate λ . Two regions are available: L with intensity C_L and H with $C_H > C_L$.

We compare two strategies (the lemma is stated for a single warm container; see the multi-container remark in §5.5): **A (warm in L)**: keep a warm container alive in L . Each arrival either reuses the warm container or — if the previous invocation was more than T_w ago — pays a cold start. **B (cold in H)**: keep no warm container. Every invocation cold-starts in H .

For Poisson inter-arrival times $X \sim \text{Exp}(\lambda)$, the probability that an arrival finds no warm container is $\Pr[X > T_w] = e^{-\lambda T_w}$, and the expected idle time of a container between consecutive uses (capped at the TTL) is $\mathbb{E}[\min(X, T_w)] = (1 - e^{-\lambda T_w})/\lambda$. Assuming $\tau_e \ll \min(1/\lambda, T_w)$:

$$E_A(\lambda) = P_a \tau_e + e^{-\lambda T_w} P_a \tau_c + \frac{1 - e^{-\lambda T_w}}{\lambda} P_i, \quad (3)$$

$$E_B = P_a(\tau_e + \tau_c). \quad (4)$$

5.2. Break-Even and the Dichotomy

Strategy A is greener than B iff $E_A \cdot C_L < E_B \cdot C_H$, i.e. $E_A/E_B < r$ where $r = C_H/C_L$. Defining $\alpha = \tau_c/\tau_e$ (cold-start-to-execution ratio) and $\beta = P_i T_w/(P_a \tau_e)$ (idle-cost ratio), and substituting $u = \lambda T_w$, the inequality $E_A < r E_B$ becomes

$$\boxed{1 + \alpha e^{-u} + \beta \frac{1 - e^{-u}}{u} < r(1 + \alpha)} \quad (5)$$

Lemma 5.1 (Cold-start carbon dichotomy). *Define $\beta_{\text{crit}}(r, \alpha) = (r - 1)(1 + \alpha)$.*

1. *If $\beta \leq \beta_{\text{crit}}$, strategy A is greener than B for all $\lambda > 0$.*
2. *If $\beta > \beta_{\text{crit}}$, there exists a unique critical arrival rate $\lambda^* > 0$ such that strategy A is greener iff $\lambda > \lambda^*$.*

Proof. Let $f(u)$ denote the LHS of (5). The limits are $f(0^+) = 1 + \alpha + \beta$ and $f(\infty) = 1$. We show f is strictly decreasing on $(0, \infty)$: differentiating, $f'(u) = -\alpha e^{-u} + \beta g'(u)$ where $g(u) = (1 - e^{-u})/u$, and a standard argument ($ue^{-u} < 1 - e^{-u}$ for $u > 0$) gives $g'(u) < 0$, hence $f'(u) < 0$. The condition $f(0^+) \leq r(1 + \alpha)$ rearranges to $\beta \leq \beta_{\text{crit}}$. Case 1: since f is decreasing and $f(0^+) \leq r(1 + \alpha)$, (5) holds for all $u > 0$. Case 2: $f(0^+) > r(1 + \alpha) > 1 = f(\infty)$; by IVT and strict monotonicity, there is a unique $u^* > 0$ with $f(u^*) = r(1 + \alpha)$, and A wins iff $u > u^*$ i.e. $\lambda > \lambda^* = u^*/T_w$. $\square$

5.3. Numerical Behaviour and Worked Example

For a webhook-handler-class function ($\tau_e = 0.3$ s, $\tau_c = 1.0$ s, $P_a = 4$ W, $P_i = 0.3$ W, $T_w = 600$ s): $\alpha \approx 3.33$, $\beta = 150$. For France/Germany ($C_L = 60$, $C_H = 350$, so $r \approx 5.83$): $\beta_{\text{crit}} \approx 20.8$. Since $\beta \gg \beta_{\text{crit}}$ a finite λ^* exists; numerical solution of (5) gives $u^* \approx 6.2$ and $\lambda^* \approx 0.010$ /s. Interpretation: a webhook function invoked more than once per ~ 100 s should keep a warm pool in France; rarer functions are greener cold-starting in Germany. For France/Poland ($r \approx 11.7$), $\lambda^* \approx 0.0048$ /s (period ~ 210 s) — the larger carbon ratio shifts break-even down by $\sim 2\times$, expanding the warm-in-L regime.

Figure 2 visualizes the break-even surface across realistic region pairs against France as the low-carbon region. The dichotomy partitions the (r, λ) plane into a “warm-in-L wins” regime (above the curve) and a “cold-in-H wins” regime (below it). Typical webhook-class arrival rates of ~ 0.5 /s sit well into the warm-in-L region for every realistic FR/X pair; typical background-class rates of $\sim 10^{-4}$ /s sit firmly in cold-in-H.

A linear approximation near the threshold, using $e^{-u} \approx 1 - u$ and $(1 - e^{-u})/u \approx 1 - u/2$, yields $u^* \approx (\beta - \beta_{\text{crit}})/(\alpha + \beta/2)$ and correspondingly $\lambda^* \approx (\beta - \beta_{\text{crit}})/[T_w(\alpha + \beta/2)]$. For our worked example this gives $\lambda^* \approx 0.0027$ /s, roughly $4\times$ smaller than the exact rate; the linearisation is a conservative lower bound on λ^* and useful when avoiding numerical root-finding.

Practical reach of the dichotomy.. The dichotomy’s structure requires one caveat for realistic FaaS function profiles: the idle-cost ratio β takes values in the range ~ 50 – 200 (for our default `webhook_handler` catalog, $\beta = 150$; longer warm-pool TTLs or lower-power idle states shift this only modestly). The break-even threshold $\beta_{\text{crit}} = (r - 1)(1 + \alpha)$ is at most ~ 47 even for extreme regional carbon ratios ($r \approx 11$, France vs. Poland). Hence for essentially all realistic FaaS deployments, $\beta > \beta_{\text{crit}}$, and the lemma’s *first* regime (warm-in-L wins for all $\lambda > 0$) is effectively unreachable: it requires either a very short warm-pool TTL ($T_w \lesssim 20$ s, well below platform defaults) or a near-50:1 inter-region carbon ratio (unphysical for grid carbon). The practical content of Lemma 5.1 for FaaS scheduling is therefore the *second* regime: the closed-form computation of λ^* that gives the scheduler a parameter-free break-even arrival rate. We retain the dichotomy statement because the first regime can arise for adjacent computing classes (long-running idle keep-alives, edge containers with $T_w \lesssim 10$ s) and because the bound $\beta \leq \beta_{\text{crit}}$ is the natural starting point for the proof, but the FaaS-specific contribution is the λ^* characterization.




Figure 2: Cold-start carbon trade-off break-even surface (Lemma 5.1). Above the curve, warming in the low-carbon region wins; below it, cold-starting in the high-carbon region wins. Annotated points are realistic France-vs-X pairs at the calibration intensities used in our worked examples ($C_L \approx 60$, $C_H \in \{200, 350, 700\}$ g CO₂eq/kWh for GB/DE/PL); real LWA 2020 mean intensities are similar in pattern (§8.3) but somewhat cleaner overall.

5.4. Idle-Energy Correction for Temporal Shifting

Lemma 5.1 addresses the *spatial* placement question. It does *not* directly address the *temporal* question: should we defer an invocation by Δt within a single region? These are distinct decisions, and conflating them is a mistake we made in our first integration attempt — with empirical consequences (§8.10, §8.13) worth reporting, because they expose a subtle aspect of the warm-pool/temporal trade-off that batch-oriented carbon-aware schedulers do not face.

Deferring an invocation by Δt extends warm-container lifetime by Δt and adds $P_i \Delta t$ of idle energy, contributing $P_i \Delta t \bar{C}_r$ carbon. For deferral to be beneficial, the savings from executing at $C_r(t + \Delta t)$ vs $C_r(t)$ must exceed this idle cost:

$$P_a \tau_e (C_r(t) - C_r(t + \Delta t)) > P_i \Delta t \bar{C}_r.$$

Dividing by $P_a \tau_e \bar{C}_r$ and defining the local carbon swing $\delta = (C_r(t) - C_r(t + \Delta t)) / \bar{C}_r$:

$$\delta > \frac{P_i \Delta t}{P_a \tau_e}.$$

The RHS is exactly the dimensionless idle-cost ratio from Lemma 5.1, with Δt in place of T_w . For our `webhook_handler` profile and a one-minute deferral, this gives $\delta > 15$, i.e. the intensity must drop by $15\bar{C}$ over the deferral — impossible. The conclusion is robust: short-horizon temporal deferral is rarely carbon-beneficial in a single region, regardless of the local diurnal swing.

Practical implication. The scheduler should not blindly pick the lowest-intensity future slot; it should charge each candidate deferral by the idle-energy cost of holding the warm container during the wait. We implement this by adding the term $P_i \Delta t \bar{C}_r / (3.6 \times 10^6)$ (grams CO₂eq) to the per-slot score in GREENFAAS’s scoring

loop. With this correction, GREENFAAS exhibits the key behaviours of §8.8–§8.13: *do-no-harm* in single-region scenarios, and stable performance across carbon-variability regimes where the ablation collapses.

5.5. Remarks and Limitations

The lemma is stated for a single warm container. The multi-container generalization changes the idle-energy term to involve the steady-state distribution of warm containers in an M/M/ n -style model. We conjecture that an analogous threshold structure holds, but the exact multi-container expression depends on the steady-state warm-pool occupancy distribution (which involves Erlang-style terms with a TTL boundary at T_w and is not closed-form), and so the precise way pool size enters the threshold remains future work; we do not claim it here. The Poisson-arrivals assumption is empirically violated by real FaaS workloads (Shahrad et al., 2020); we test the lemma’s qualitative prescription on the schema-faithful bursty workload calibrated to the Azure 2021 trace in §8.2, and on real LWA carbon data in §8.3, and find it robust in both. PUE heterogeneity extends the analysis by replacing r with $r' = (C_H \cdot \text{PUE}_H)/(C_L \cdot \text{PUE}_L)$.

6. The GreenFaaS Algorithm

This section presents the GREENFAAS scheduler. §6.1 gives a high-level narrative. §6.2 presents the algorithm in pseudocode. §6.3 describes the SLA-tiered policy. §6.4 analyses time and space complexity. §6.5 discusses arrival-rate estimation and jitter.

6.1. Design Overview

GREENFAAS schedules each invocation independently at its arrival, producing a decision $(r(i), s(i))$. The design has three layers: **(1)** an arrival-rate estimator tracks per-(function, home-region) arrival rates online with an EWMA; **(2)** a *region gate* applies Lemma 5.1 to each candidate region given the current $\hat{\lambda}$ and present carbon intensities, dropping regions that fail the gate not because they could not save carbon *now* but because routing this function class to them on a sustained basis would lose carbon to warm-pool idle energy; **(3)** a *per-slot scorer* enumerates a small grid of (region, start-time) pairs over the admitted regions and the SLA-bounded deferral window, selecting the pair with minimum expected carbon, with an idle-energy term that charges deferred invocations for the additional warm-container holding cost (§5.4).

These three layers are stateless across invocations except for the arrival-rate estimator. There is no global queue, no inter-invocation coordination, no pile-up risk from synchronised defer slots (we use a per-function deterministic jitter to ensure this; §6.5). The scheduler can run in parallel across invocations with no locks beyond the rate update.

6.2. Algorithm

We present GREENFAAS as Algorithm 1 with sub-procedures LEMMAGATE and SCORESLOTS. We write $\mathcal{S}(t)$ for the platform state at t : warm-pool counts $W_{r,f}(t)$ and in-flight counts $N_r(t)$.

The three energy quantities are $E_{\text{exec}}(f, r) = P_a(f)\tau_e(f) \text{PUE}_r/(3.6 \times 10^6)$, $E_{\text{cs}}(f, r) = P_a(f)\tau_c(f) \text{PUE}_r/(3.6 \times 10^6)$, and $E_{\text{idle}}(f, r, \Delta t) = P_i(f)\Delta t \text{PUE}_r/(3.6 \times 10^6)$, all in kWh; multiplying by the appropriate carbon intensity gives gCO₂eq. The use of $\bar{C}_r$ for the idle term (rather than $\hat{C}_r(t_k)$) is the §5.4 correction: idle energy accrues *over* the deferral window, not at a single instant.

Operational semantics of the deferral idle term. GreenFaaS does *not* reserve a dedicated warm container during a deferral: a deferred invocation ($k > 0$) is conservatively treated as a cold start (line 8, cold $\leftarrow \neg(w_0 \wedge k=0)$), because the platform makes no guarantee that a warm container will survive the wait under contention. The $E_{\text{idle}}(f, r, t_k - t_0) \bar{C}_r$ term is therefore not a literal per-invocation energy charge for a reserved container; it is a *conservative opportunity-cost penalty* that approximates the additional warm-pool exposure the deferral induces at the pool level (a delayed invocation lengthens the interval over which the pool’s warm capacity is held without serving work). This conservatism is deliberate: it biases the scorer against deferral unless the carbon saving clearly outweighs the penalty, which is precisely the mechanism that yields the do-no-harm property (§8.11). An alternative model that reserves a warm container during

Algorithm 1 GREENFAAS($i, \mathcal{S}, \hat{C}$)

```
1:  $\hat{\lambda} \leftarrow \text{UPDATERATE}(f, h, a(i))$ 
2: if  $\ell(f) = \text{INTERACTIVE}$  then return  $(h, a(i))$ 
3: end if
4:  $(\rho^*, \Delta^*) \leftarrow \text{CLASSLIMITS}(\ell(f))$ 
5:  $\mathcal{R}_0 \leftarrow \{r : \rho(h, r) \leq \rho^* \wedge N_r/K_r < u_{\max}\}$ 
6:  $\mathcal{R}^* \leftarrow \{r \in \mathcal{R}_0 : \text{LEMMAGATE}(f, h, r, \hat{\lambda}, \dots)\}$ 
7: if  $\mathcal{R}^* = \emptyset$  then
8:    $\mathcal{R}^* \leftarrow \{h\}$ 
9: end if
10: return SCORESLOTS( $f, h, \mathcal{R}^*, \Delta^*, \hat{C}, \mathcal{S}, a(i)$ )
```

Algorithm 2 LEMMAGATE($f, h, r, \hat{\lambda}, \mathcal{S}, \hat{C}, t$)

```
1: if  $r = h$  then return true
2: end if
3:  $C_r \leftarrow \hat{C}_r(t), C_h \leftarrow \hat{C}_h(t)$ 
4: if  $C_r \geq C_h$  then return true ▷ admit; scorer compares instantaneous carbon
5: end if
6:  $\alpha \leftarrow \tau_c(f)/\tau_e(f)$ 
7:  $\beta \leftarrow P_i(f)T_w/(P_a(f)\tau_e(f))$ 
8:  $\text{ratio} \leftarrow C_h/C_r$ 
9:  $\beta_{\text{crit}} \leftarrow (\text{ratio} - 1)(1 + \alpha)$ 
10: if  $\beta \leq \beta_{\text{crit}}$  then return true ▷ Regime 1
11: end if
12:  $\lambda^* \leftarrow \text{SOLVEBREAK EVEN}(\text{ratio}, \alpha, \beta, T_w)$ 
13: return  $\hat{\lambda} > \lambda^*$ 
```

deferral — charging literal idle energy and setting cold $\leftarrow \neg(w_0 \wedge k\delta < T_w)$ — would relax this conservatism; we adopt the no-reservation model because it matches the default behaviour of FaaS platforms that do not pin containers to pending requests.

SOLVEBREAK EVEN solves (5) for $u^* > 0$ by bisection on the strictly decreasing LHS, returning $\lambda^* = u^*/T_w$. Convergence to 10^{-9} requires ≤ 200 iterations in practice.

6.3. SLA-Tiered Policy

The class limits (ρ^*, Δ^*) in line 3 of Algorithm 1 encode SLA tiering. Defaults:

Class	ρ^* (RTT)	Δ^* (defer)
INTERACTIVE	0 ms	0 s
DEFERRABLE	80 ms	$\min(60 \text{ s}, D - \tau_e)$
BACKGROUND	∞	$D - \tau_e$

DEFERRABLE invocations are capped at 60s of deferral even when their deadline is longer, which keeps the temporal-shift horizon below the warm-pool TTL ($T_w = 600 \text{ s}$ by default) and ensures the idle-energy term remains a small correction rather than the dominant cost. BACKGROUND invocations use the full deadline.

6.4. Complexity

Let $R = |\mathcal{R}|$, N the number of forecast slots per region, B the bisection iterations in SOLVEBREAK EVEN. UPDATERATE is $O(1)$. The candidate filtering is $O(R)$. The region gate calls LEMMAGATE once per

Algorithm 3 SCORESLOTS($f, h, \mathcal{R}^*, \Delta^*, \hat{C}, \mathcal{S}, t_0$)

```
1:  $(r^*, s^*, S^*) \leftarrow (\perp, \perp, +\infty)$ 
2: for each  $r \in \mathcal{R}^*$  do
3:    $\delta \leftarrow$  forecast step;  $N \leftarrow \lceil \Delta^* / \delta \rceil + 1$ 
4:    $\bar{C}_r \leftarrow$  long-run mean of  $C_r$ 
5:    $w_0 \leftarrow \mathbf{1}[W_{r,f}(t_0) > 0]$ 
6:   for  $k = 0, 1, \dots, N - 1$  do
7:      $t_k \leftarrow t_0 + k\delta + \text{jitter}(f, i, k) \cdot \delta$ 
8:      $\text{cold} \leftarrow \neg(w_0 \wedge k = 0)$ 
9:      $\text{eta} \leftarrow (t_k - t_0) + \rho(h, r) + \text{cold} \cdot \tau_c(f) + \tau_e(f)$ 
10:    if  $\text{eta} > D(f)$  then continue
11:    end if
12:     $C_k \leftarrow \hat{C}_r(t_k)$ 
13:     $S \leftarrow E_{\text{exec}}(f, r)C_k + \text{cold} \cdot E_{\text{cs}}(f, r)C_k + E_{\text{idle}}(f, r, t_k - t_0)\bar{C}_r$ 
14:    if  $S < S^*$  then  $(r^*, s^*, S^*) \leftarrow (r, t_k, S)$ 
15:    end if
16:  end for
17: end for
18: if  $r^* = \perp$  then return  $(h, t_0)$ 
19: end if
20: return  $(r^*, s^*)$ 
```

candidate, each $O(B)$. SCORESLOTS enumerates $O(R \cdot N)$ slots, each $O(1)$ given precomputed $\bar{C}_r$. **Per-invocation time:** $O(R \cdot (B + N))$. For default settings ($R \leq 9$, $N \leq 12$ for DEFERRABLE, $N \approx 720$ for BACKGROUND), this is well below $100 \mu\text{s}$ per invocation in practice. **Per-invocation space:** $O(R)$. The arrival-rate estimator uses $O(|\mathcal{F}| \cdot R)$ state amortised across time.

6.5. Arrival-Rate Estimation and Jitter

UPDATERATE maintains a per-(function, region) EWMA: $\hat{\lambda} \leftarrow \gamma \cdot (t - t_{\text{prev}})^{-1} + (1 - \gamma)\hat{\lambda}_{\text{prev}}$ with smoothing coefficient $\gamma = 0.2$ (we use γ here to avoid collision with $\alpha = \tau_c/\tau_e$ from §5). Before the first update, $\hat{\lambda}$ is initialised to a small default (0.01/s) making the gate conservative for cold-start functions.

Without jitter, all invocations of the same function in the same deferral window would target the same defer slot — exactly the Wait-Awhile failure mode of synchronised pile-up. We apply a deterministic per-function offset $\text{jitter}(f, i, k) = (H(f \oplus i) \bmod M) / M \cdot \phi$ for $k \geq 1$, where H is a hash, $M = 65535$, and $\phi = 0.25$ is the max fraction of a slot width by which jitter shifts the candidate time. The jitter is zero at $k = 0$ so that immediate execution remains a candidate.

Forecasts.. $\hat{C}$ may be the true carbon trace (**perfect**), the true trace within a cutoff with Gaussian-noise decay beyond (1h, 24h), or a constant equal to each region’s historical mean (**none**). Forecast accuracy is a sensitivity axis in §8.12; we will see that GREENFAAS is essentially forecast-invariant, an unintended but welcome consequence of the design.

7. Simulator Implementation

We implemented GREENFAAS-sim, a single-process discrete-event simulator in Python, alongside the scheduler. The design follows the abstractions of **faas-sim** (function-level FaaS simulation) and **vessim** (carbon co-simulation) (Edgerun, 2022; dos-group, 2023), but is a clean-room implementation rather than a fork — written from scratch to give us tight control over event ordering, warm-pool accounting, and the extension points for trace loaders. The full implementation, including all baseline schedulers, is approximately 1,700 lines of substantive Python with no required dependencies beyond the standard library.

7.1. Event Model

The simulator maintains a min-heap of timestamped events of four kinds. **ARRIVAL**: a new invocation enters the system; the scheduler is consulted exactly once and the decision committed for the lifetime of the invocation. **START**: the chosen start time of a previously-scheduled invocation; execution begins, energy and carbon are charged. **COMPLETION**: an invocation finishes; the container becomes warm-and-idle for up to T_w seconds. **WARMEXPIRE**: a warm container reaches its TTL without being reused and is evicted; accumulated idle energy is finalised.

The decoupling of **ARRIVAL** from **START** is what makes temporal deferral possible without special-casing in the event loop: the scheduler returns a `start_time` that may be in the future, and the simulator simply enqueues the **START** event at that time. Reusing a warm container is signalled by the scheduler’s `use_warm` flag and realised at the **START** event by decrementing the warm-pool count.

7.2. Warm-Pool Accounting

For each (r, f) pair, the simulator tracks $W_{r,f}(t)$ and a queue of pending **WARMEXPIRE** events. A container becomes warm at **COMPLETION** with scheduled expiry T_w seconds later. If reused before expiry, the warm count is decremented and the corresponding **WARMEXPIRE** is matched at the **START** event; the container’s lifetime ($t_{\text{end}} - t_{\text{start}}$) is logged. If expiry fires first, the **WARMEXPIRE** event performs eviction and logs the full T_w lifetime.

At simulation end, total warm-container lifetime per (r, f) pair is multiplied by $P_i(f)$, PUE_r , and the region’s *mean* carbon intensity $\bar{C}_r$. As discussed in §4.1, this mean-intensity term is an approximation of the exact interval integral $\int W_{r,f}(t)C_r(t)dt$; it is exact when idle time is distributed uniformly over the horizon and conservative otherwise. The simulator therefore reports results under the mean-intensity approximation justified in §4.1, matching the convention used in GREENFAAS’s per-slot scoring (Algorithm 3, line 13).

7.3. Extension Points

Three abstractions cleanly separate the algorithm from the data sources. **CarbonModel** provides `intensity(region, t)` and `forecast(region, t, horizon, accuracy)`; the default implementation reads from per-region **CarbonTrace** objects with arbitrary step sizes, and real LWA / ENTSO-E format CSVs (dos-group, 2021) are loaded via `load_carbon_csv` into the same type. The scheduler does not know whether the carbon data is synthetic or real. **Workload generators** produce `List[Invocation]`; the default emits a non-homogeneous Poisson stream with Zipf popularity, and the real-trace loaders emit the same type from published Azure schemas. **Schedulers** implement a single method `schedule(invocation, state, carbon) -> ScheduleDecision`; all five schedulers in our evaluation conform to this interface and are swappable.

7.4. Validation

The simulator’s correctness rests on three checks. (1) *Conservation*: total reported execution energy equals $\sum_i P_a(f(i))\tau(i)$ within floating-point precision; total warm-idle energy equals $\sum_{(r,f)} P_i(f) \cdot \text{lifetime}_{r,f}$. Verified by direct comparison with hand-computed values on small synthetic workloads. (2) *Schedule determinism*: re-running a fixed scheduler on a fixed invocation stream produces identical decisions; verified for all five schedulers. (3) *Lemma-implementation agreement*: the closed-form Lemma 5.1 and the per-invocation carbon computed by direct simulation agree at 56 grid points spanning $r \in [1.1, 11.7]$ and $\lambda \in [10^{-4}, 1]/s$, with *zero* mismatches: the brute-force simulator confirms the lemma’s predicted regime boundary (warm-in-low-carbon versus cold-in-high-carbon) at every grid point, and confirms the predicted break-even rate λ^* within the grid resolution (`scripts/verify_tradeoff.py`). The third check in particular gives us confidence that the scheduler’s lemma-driven decisions and the simulator’s energy accounting are mutually consistent — addressing a concern that arises because the lemma was derived analytically and the simulator was implemented independently.

8. Evaluation

This section presents our experimental evaluation of GREENFAAS against four carbon-aware baselines. §8.1 describes the trace data and reconstruction methodology. §8.2 reports headline results. §8.8–§8.14 sweep five sensitivity axes (topology, SLA mix, forecast accuracy, carbon variability, workload intensity). §8.15 synthesizes the cross-axis findings.

8.1. Methodology

Workload traces. Our evaluation uses two kinds of workload. The controlled sensitivity sweeps (§8.10–§8.14) are driven by a schema-faithful *synthetic* workload, because their purpose is to vary individual parameters (SLA mix, forecast accuracy, carbon variability, arrival rate) cleanly across wide ranges. The fully-real validation (§8.6) is driven by the real *Azure Functions Invocation Trace 2021* (Zhang et al., 2021; Microsoft, 2021), a per-invocation trace with schema (`app`, `func`, `end_timestamp`, `duration`); the 2021 loader converts each row to an `Invocation` by computing `arrival_time = end_timestamp - duration` and mapping (`app`, `func`) pairs to stable function IDs. Our loaders also accept the *Azure Functions 2019* dataset (Shahrad et al., 2020), which is *aggregated* (per-minute invocation counts plus per-function duration percentiles); the 2019 loader reconstructs a per-invocation stream by drawing arrival times uniformly within each minute bin and sampling per-invocation durations from the empirical percentile distribution via linear inverse-CDF interpolation. Both loaders accept the exact published schemas, so the same evaluation code runs end-to-end on either real or synthetic input.

Synthetic workload calibration. The synthetic workload used in the sensitivity sweeps is calibrated to the Shahrad et al. (2020) and Zhang et al. (2021) characterizations: non-homogeneous Poisson arrivals with diurnal shape, Zipf popularity with $s = 1.2$, heavy-tailed durations, and the SLA mix described below. We use synthetic input for the sweeps because controlled parameter variation across multiple axes requires generating workloads the real trace does not contain (e.g. a 25% shiftable fraction, or a forecast-error level of 0.3); the fully-real experiment of §8.6 confirms that the headline ordering established on synthetic data reproduces on the real Azure 2021 trace paired with time-aligned real carbon.

Latency-class assignment. Neither Azure dataset explicitly labels SLA tiers. For the 2019 trace the loader maps the explicit `Trigger` field: HTTP $\rightarrow$ INTERACTIVE; queue, event-grid, orchestration $\rightarrow$ DEFERRABLE; timer, storage $\rightarrow$ BACKGROUND. For the 2021 trace (no trigger field) the loader classifies by observed mean duration: $< 1\text{ s} \rightarrow$ INTERACTIVE (1s SLA); $1\text{--}30\text{ s} \rightarrow$ DEFERRABLE (60s SLA); $> 30\text{ s} \rightarrow$ BACKGROUND (1 hour SLA). The synthetic workload uses the same class distribution as the 2019 trigger-field histogram.

Carbon traces. We use the publicly-released hourly carbon-intensity traces from Wiesner et al. (2021); dos-group (2021) (DE, GB, FR, US-CAISO, 2020), computed from ENTSO-E and CAISO generation-mix data using IPCC factors. For experiments requiring a coal-belt region (§8.9) we extend with a calibrated PL trace whose annual mean (791 g CO₂eq/kWh) matches the value LWA documents for Poland in 2020 and whose coal-base-load diurnal shape uses the same generation factors as the LWA computation pipeline; we label it “PL (synth., LWA methodology)” wherever it appears. The simulator resamples all carbon traces to a uniform 5-minute resolution via linear interpolation.

Baselines. **FIFO** (carbon-unaware lower bound). **Wait-Awhile** (Wiesner et al., 2021): temporal deferral with threshold $C^* = 200\text{ g CO}_2\text{eq/kWh}$. Two variants are reported: *Wait-Awhile-Batch* (`max_defer` = 1800s, Wiesner et al.’s batch-job default) and *Wait-Awhile-FaaS* (`max_defer` = 60s, matching GREENFAAS’s DEFERRABLE horizon). The two variants together address an obvious reviewer concern (§8.4). **Spatial**: route to the lowest-current-carbon region within a class-dependent RTT budget. **GreenFaaS-v1**: the GREENFAAS scheduler *without* the §5.4 idle-energy correction, included as an ablation.

Table 1: Synthetic 24h workload, 5 regions, 173k invocations, **mean $\pm$ std across 5 seeds** (42–46). Wait-Awhile-Batch uses the original 1800 s defer horizon; see §8.4 for the fair-parameterization variant.

Scheduler	Carbon (g)	vs FIFO
FIFO	51.57 ± 0.37	0.0%
Wait-Awhile-Batch	209.66 ± 2.51	$-306.5\% \pm 4.3\%$
Spatial	13.97 ± 0.07	$+72.91\% \pm 0.26\%$
GreenFaaS-v1	14.73 ± 0.13	$+71.43\% \pm 0.40\%$
GREENFAAS	14.52 ± 0.07	$+71.84\% \pm 0.32\%$

8.2. Headline Results

We report headline results on a canonical synthetic experiment and validate them on real carbon data in §8.3. The canonical workload is 24h long with 173k invocations across 5 regions including France as a low-carbon refuge and Poland at the high-carbon extreme.

Three observations frame the rest of §8. *First*, the batch-oriented Wait-Awhile scheduler with its default 1800 s deferral horizon *more than triples* operational carbon ($-306.5\% \pm 4.3\%$ on synthetic, $-176.2\% \pm 2.5\%$ on real LWA across 5 seeds), an empirical confirmation that batch carbon-aware techniques do not transfer to FaaS. §8.4 shows that shorter, FaaS-appropriate deferral horizons (60 s, 30 s) do not rescue Wait-Awhile — they cause it to degenerate to a FIFO no-op, because its absolute-threshold criterion is rarely met inside such windows. We attribute the long-horizon failure mode in §8.13 to the synchronised-defer pile-up effect noted in §5.4. *Second*, GREENFAAS reaches $71.84\% \pm 0.32\%$ reduction relative to FIFO, within 1.1 percentage points of pure spatial routing ($72.91\% \pm 0.26\%$); the GreenFaaS–Spatial gap is small but statistically detectable (see the significance paragraph below), with Spatial consistently edging GREENFAAS. The two schedulers are indistinguishable on SLA violation rate, p95 latency, and cold-start rate. *Third*, the ablation between GreenFaaS-v1 and GREENFAAS shows a meaningful but modest gap on this canonical workload ($71.43\% \pm 0.40\%$ vs $71.84\% \pm 0.32\%$). The interesting separation appears not on canonical workloads but in the sensitivity sweeps of §8.10 and §8.13, and in the coal-belt topology (§8.9) where the ablation’s variance blows up by an order of magnitude (std jumps from 0.4% to 9.8%) — v1 produces unpredictably bad outcomes when the spatial axis offers limited refuge.

The real-carbon validation in §8.3 reproduces the same ordering on real LWA 2020 ENTSO-E / CAISO data, with magnitudes shifted modestly (GreenFaaS $64.63\% \pm 0.24\%$ reduction vs $\sim 72\%$ on synthetic) because real 2020 European carbon was cleaner overall than our synthetic calibration assumed.

Statistical significance of the key gaps. We report paired statistical tests (paired t -test and Wilcoxon signed-rank) for the load-bearing scheduler comparisons, computed over the per-seed (synthetic, $n = 5$) and per-window (real window-shift, $n = 5$) carbon measurements (`scripts/stat_tests.py`). The GreenFaaS–Spatial gap is small but statistically detectable: on synthetic data the 0.55 g mean difference is significant (paired t , $p = 0.0003$), and on the real window-shift data the 0.11 g difference is significant ($p = 0.015$). In both cases Spatial is reliably the marginally better of the two; the practical magnitude (sub-percentage-point) is what licenses the claim that GREENFAAS *matches* Spatial, not that the gap is zero. The GreenFaaS-vs-v1 ablation gap is significant on both synthetic ($p = 0.035$) and real ($p = 0.003$) data, confirming the §5.4 idle-energy correction has a real effect. The Wait-Awhile and Lechowicz losses to FIFO are significant on synthetic data; on the real window-shift data with $n = 5$, the Lechowicz loss is significant ($p = 0.042$) while the Wait-Awhile loss does not reach significance ($p = 0.089$), reflecting the smaller real-data effect size and the limited window count. We note that the Wilcoxon signed-rank test cannot reach $p < 0.05$ with $n = 5$ (its minimum two-sided p -value is 0.0625), so we rely on the paired t -test for significance at $n = 5$ and report Wilcoxon as a non-parametric robustness check; a larger n would be needed to discriminate the smallest effects non-parametrically. Given $n = 5$, we treat the *practical effect size* — the sub-percentage-point GreenFaaS–Spatial gap and the multi-point batch-scheduler losses — rather than the p -values as the load-bearing evidence; the tests are reported to characterize the gaps, not to rest the paper’s claims on significance.

at small n . The five seeds (and the five windows) are independent draws; we apply no multiple-comparison correction and do not claim any single p -value as decisive.

8.3. Real Carbon-Intensity Validation

To validate that the headline findings of §8.2 are not artefacts of synthetic carbon traces, we re-ran the canonical 24-hour experiment substituting the synthetic carbon model with the real ENTSO-E / CAISO 15-minute-resolution carbon-intensity data released by Wiesner et al. (2021); dos-group (2021). The workload in *this* experiment is the schema-correct synthetic stream, holding the workload fixed while varying only the carbon model; the fully-real version that also substitutes the real Azure 2021 trace is reported in §8.6.

Carbon profile. The real LWA traces over our 24-hour window exhibit mean intensities of 39.9 g/kWh (France), 235.2 g (Germany), 234.6 g (Great Britain), and 319.0 g (US-CAISO). These are approximately 15–25% lower than the calibration values we used in the synthetic experiments (60, 350, 220, 250 g respectively), reflecting ENTSO-E’s actual 2020 generation mix rather than published annual means. The diurnal swing is real and non-trivial: GB ranges from 121 to 336 g/kWh on this day, a $\sim 3\times$ ratio; CAISO from 243 to 352, a $\sim 1.5\times$ ratio. France is essentially flat (38–43 g) due to its nuclear-dominated mix.

Results. Table 2 reports the same five schedulers on real carbon data; compare to Table 1. The qualitative findings of §8.2 hold without modification.

Table 2: Real LWA carbon, synthetic workload (24h, 173k invocations, 4 regions: DE, FR, GB, US-CAISO), **mean $\pm$ std across 5 seeds**. Compare to Table 1.

Scheduler	Carbon (g)	vs FIFO
FIFO	33.34 ± 0.24	0.0%
Wait-Awhile	92.08 ± 0.36	$-176.2\% \pm 2.5\%$
Spatial	11.63 ± 0.01	$+65.12\% \pm 0.23\%$
GreenFaaS-v1	14.40 ± 0.17	$+56.79\% \pm 0.56\%$
GREENFAAS	11.79 ± 0.03	$+64.63\% \pm 0.24\%$

Three observations. *First*, the Wait-Awhile failure mode reproduces on real data ($-176\% \pm 2.5\%$ over FIFO, vs $-306\% \pm 4.3\%$ on synthetic). The mechanism is the same (synchronised-defer pile-up inflating cold-start and idle-energy), and the magnitude reduction is consistent with the smaller diurnal swing in the real data. *Second*, the GREENFAAS–Spatial gap is 0.49 ± 0.33 pp, right on the edge of statistical significance; Spatial slightly edges GREENFAAS. We attribute this to the absence of a coal-belt region in the LWA dataset; under §8.8, the coal-belt scenario is where GREENFAAS pulls ahead, and the four-region LWA mix (DE/FR/GB/CAISO) lacks such a region. *Third*, GreenFaaS-v1’s 7pp gap below GREENFAAS on real data validates the §5.4 idle-energy correction in the real-carbon regime, with the same cold-start-rate increase (0.34% vs 0.03%) seen on synthetic data.

In summary, real carbon data does *not* change any qualitative finding of §8.2. The magnitude of GREENFAAS’s reduction vs FIFO is somewhat smaller on real data (64% vs 72%) because the real European 2020 mix is cleaner overall than our synthetic calibration assumed. The accompanying topology sweep on real LWA 2020 traces extended with a calibrated PL trace (§8.9, Table 8) extends this validation across topology variations and tempers the §8.8 coal-belt claim. The fully-real validation — combining the real Azure 2021 per-invocation trace with time-aligned ElectricityMaps carbon data — is reported in §8.6, and confirms the qualitative findings reported here on real workload as well as real carbon.

8.4. Wait-Awhile under FaaS-Appropriate Parameters

The headline Table 1 uses Wait-Awhile with its original 1800s deferral horizon (Wiesner et al.’s batch-job default (Wiesner et al., 2021)). A natural reviewer concern is that 1800s is adversarial against FaaS functions with sub-minute SLA deadlines. We address this directly by reporting Wait-Awhile with shorter deferral horizons matched to FaaS-appropriate timescales.

Table 3: Wait-Awhile under three deferral horizons. Synthetic 24h workload; real-LWA-carbon results are nearly identical.

Scheduler	Carbon (g)	vs FIFO	Cold rate
FIFO	51.0	0.0%	0.07%
Wait-Awhile (defer=1800 s)	206.6	−305.0%	2.02%
Wait-Awhile (defer=60 s)	51.0	0.0%	0.07%
Wait-Awhile (defer=30 s)	51.0	0.0%	0.07%
Spatial	14.0	+72.6%	0.03%
GREENFAAS	14.6	+71.4%	0.05%

The empirical finding is more nuanced than the reviewer concern anticipated. With a FaaS-appropriate deferral horizon (60 s or 30 s), *Wait-Awhile becomes bit-identical to FIFO*: no future slot meets its absolute-threshold criterion within the short window, so it always falls back to immediate execution. The batch-default horizon (1800 s) produces catastrophic carbon inflation; the FaaS-appropriate horizons produce no carbon decisions at all. Wait-Awhile’s batch-job design has no parameter sweet spot for FaaS arrival rates and SLA deadlines.

This finding refines the headline results of §8.2. The statement “Wait-Awhile more than doubles operational carbon on FaaS” is true under Wait-Awhile’s default parameters; under FaaS-appropriate parameters the failure mode is different (the scheduler is starved of valid deferral options and degenerates to FIFO). Either way, Wait-Awhile as designed does not produce carbon savings on FaaS workloads, but the mechanism of failure depends on the parameterization: catastrophic at long horizons, no-op at short horizons. GREENFAAS’s SLA-tiered deferrable horizon (60 s for DEFERRABLE class) plus relative carbon-aware scoring (rather than absolute thresholds) is precisely what avoids both failure modes.

Threshold parameterization. The experiments above vary Wait-Awhile’s deferral *horizon* (1800 s vs 60 s) while holding its absolute carbon threshold fixed at $C^* = 200 \text{ gCO}_2/\text{kWh}$. A lower threshold would defer fewer invocations (approaching FIFO sooner) and a higher one would defer more (amplifying the catastrophic mode at long horizons); neither escapes the underlying problem, which is that an *absolute* carbon threshold combined with any fixed horizon cannot adapt to the warm-pool idle-energy coupling that dominates FaaS. We fix C^* at a value near the midpoint of our regions’ intensity range so that the comparison isolates the horizon effect; a full two-dimensional (horizon $\times$ threshold) sweep is a natural extension but would not change the qualitative conclusion, since the failure is structural (absolute thresholding) rather than a matter of threshold tuning.

8.5. Comparison with the Provable Lechowicz Algorithm

A reviewer concern (and a natural one) is that our headline comparison is against three heuristics (FIFO, Wait-Awhile, Spatial) rather than against a provably-optimal online algorithm. We address this by implementing a FaaS-adapted port of the one-way trading algorithm of Lechowicz et al. (2023), which provides an optimal competitive ratio $\sqrt{U/L}$ for deadline-constrained carbon-aware scheduling in a simpler model. Our port: for each invocation, examine the carbon forecast over the SLA-deadline window, compute the threshold $\Phi^* = \sqrt{U \cdot L}$ (where U, L are the maximum and minimum forecast intensities), commit at the first slot below Φ^* , otherwise commit at the deadline.

The finding is striking: **the provably-optimal-competitive-ratio algorithm performs worse than the heuristic Wait-Awhile on real carbon data, and worst of all on the single-region setup where its theory applies cleanest.** (On *synthetic* carbon the two batch schedulers rank the other way — Lechowicz −238.4% vs Wait-Awhile −306.5%, so Lechowicz is marginally the less catastrophic — but on real carbon the ordering reverses, Lechowicz −336.6% vs Wait-Awhile −176.2%; either way both fail badly.) In the most-favourable-to-Lechowicz setup (real single-region DE, no spatial axis), GREENFAAS and Spatial correctly identify zero opportunity and match FIFO byte-for-byte, while Lechowicz inflates carbon by 320% above FIFO.

Table 4: Lechowicz et al. (Lechowicz et al., 2023) one-way trading algorithm vs. our baselines, 5-seed mean $\pm$ std.

Setup	Scheduler	vs FIFO	Cold-start
Synthetic 5-region	Wait-Awhile	$-306.5\% \pm 4.3\%$	2.03%
	Lechowicz	$-238.4\% \pm 3.2\%$	2.39%
	Spatial	$+72.9\% \pm 0.3\%$	0.03%
	GREENFAAS	$+71.8\% \pm 0.3\%$	0.05%
Real LWA 4-region	Wait-Awhile	$-176.2\% \pm 2.5\%$	1.29%
	Lechowicz	$-\mathbf{336.6\%} \pm \mathbf{4.3\%}$	2.84%
	Spatial	$+65.1\% \pm 0.2\%$	0.03%
	GREENFAAS	$+64.6\% \pm 0.2\%$	0.03%
Real single-region DE	Wait-Awhile	$-176.7\% \pm 6.3\%$	1.61%
	Lechowicz	$-\mathbf{320.2\%} \pm \mathbf{2.9\%}$	2.44%
	Spatial	$0.0\% \pm 0.0\%$	0.02%
	GREENFAAS	$0.0\% \pm 0.0\%$	0.02%

The mechanism is exactly the §5.4 contribution. Lechowicz’s competitive analysis assumes deferral is free of side effects: the cost of holding a job pending until commit time is treated as zero. For FaaS, this assumption is violated by warm-pool idle energy: deferring an invocation extends the warm container’s idle period, charging $P_i \Delta t \bar{C}_r$ in carbon that the competitive analysis does not see. The threshold $\Phi^* = \sqrt{U \cdot L}$ aggressively defers any invocation arriving above Φ^* , even when the warm container’s idle-energy cost will exceed the execution savings — which is, given typical FaaS values of $\beta = P_i T_w / (P_a \tau_e) \in [50, 200]$, almost always. The cold-start rate of 2.4–2.8% (vs. GREENFAAS’s 0.02–0.05%) is the direct evidence: the algorithm aggressively defers past warm-pool TTLs and pays cold-start penalties on most invocations.

This is, in our reading, the strongest empirical evidence for the paper’s central technical claim. A provably-optimal algorithm for deadline-constrained one-way trading, when ported to FaaS scheduling without accounting for warm-pool idle-energy coupling, becomes the worst-performing scheduler in our study. The closed-form idle-energy correction of §5.4 is precisely what makes the difference; without it, provability does not rescue the algorithm. §9 formalizes the warm-pool-coupled cost model and states Conjecture 9.2: that Lechowicz-style algorithms have $\Omega(\beta)$ competitive ratio in a multi-container regime that captures FaaS production behavior. We have not proven this conjecture; the single-container model does not admit the natural N-invocation construction (the idle energy of one shared container cannot be charged per pending invocation without double-counting, see §9.3), and the empirical evidence here is consistent with the conjecture but does not establish it.

Honest caveats. (i) Our port follows Lechowicz et al.’s algorithmic prescription faithfully but does not extend their competitive-ratio proof to the warm-pool-coupled cost. The non-separable cost structure of GREENFAAS (§4.2) is outside the standard one-way trading framework. (ii) The Lechowicz framework was designed for batch jobs with hours-long deadlines, similar to Wait-Awhile’s intended regime; we are not surprised that a FaaS-adapted port fails, and the paper’s finding is that *neither* batch-oriented framework (heuristic or provably-optimal) transfers to FaaS without addressing the idle-energy coupling. A FaaS-native online algorithm with competitive-ratio guarantees over the warm-pool-coupled problem remains the most interesting theoretical extension this work suggests. (iii) Our port computes the one-way trading bounds U and L from the carbon-intensity forecast over each invocation’s deadline window, whereas the original algorithm assumes globally-known a-priori bounds. Using deadline-local bounds adapts the threshold $\Phi^* = \sqrt{UL}$ to each invocation’s forecast window; this is a necessary adaptation to the non-stationary carbon setting, and it may make the threshold more or less aggressive than the global-bound version depending on the local carbon range.

8.6. Time-Aligned Fully-Real Validation

The strongest empirical validation in the paper combines the real Azure Functions 2021 per-invocation trace with real ElectricityMaps carbon-intensity data from the *same calendar window* as the trace (12–25 January 2021). All other experiments in §8 use synthetic-but-schema-faithful workload calibrated to Azure characterizations (Shahrad et al., 2020; Zhang et al., 2021; Microsoft, 2021), or real carbon with synthetic workload, or real workload with carbon from a later year. This section substitutes both axes with real time-aligned data and reports the findings.

Data sources..

- **Workload:** the `AzureFunctionsInvocationTraceForTwoWeeksJan2021` trace from Zhang et al. (2021); Microsoft (2021) (1.98M invocations over 14 days, loaded as a 24h slice giving 80,990 invocations across 119 unique functions; round-robin assignment to the 5 regions). The trace provides exact per-invocation arrival timestamps and execution durations.
- **Carbon:** ElectricityMaps hourly carbon-intensity feeds for the five regions over 12–25 January 2021 (DE 396.6 g mean, FR 89.4 g, GB 293.3 g, US-CAISO 319.6 g, PL 801.1 g; range from FR’s diurnal flat 73–108 g to PL’s coal-base-load 690–869 g, giving a carbon-ratio span of roughly 11× across the topology).

Table 5: Real Azure 2021 trace + real ElectricityMaps Jan 2021 carbon data, 24h, 80,990 invocations, 5 regions.

Scheduler	Carbon (g)	vs FIFO	SLA	Cold
FIFO	231.33	0.0%	0.39%	1.68%
Wait-Awhile-Batch	239.29	−3.44%	0.39%	2.25%
Lechowicz	242.91	−5.00%	0.39%	2.50%
Spatial	73.11	+ 68.39%	0.39%	1.68%
GreenFaaS-v1	78.33	+66.14%	0.64%	2.99%
GREENFAAS	73.22	+ 68.35%	0.39%	1.68%

Headline result on fully-real time-aligned data.. Four observations.

(1) **GreenFaaS–Spatial gap collapses to 0.04 percentage points** on fully-real time-aligned data — tighter than the synthetic 1.1pp gap and tighter than the real-LWA-carbon 0.5pp gap. On real workload + real carbon, GreenFaaS and Spatial are empirically indistinguishable. The *topology-robustness* claim from §8.8 stands: GreenFaaS matches Spatial when Spatial works, and as §8.9 shows, falls back to FIFO byte-for-byte when Spatial does not.

(2) **Wait-Awhile and Lechowicz failure modes are *quantitatively softer* on real workload than on synthetic.** Wait-Awhile loses only 3.44% on real workload + real carbon, vs −306% on synthetic. Lechowicz loses 5.00%, vs −238% on synthetic. The qualitative finding (batch carbon-aware techniques fail on FaaS) holds; the catastrophic magnitudes seen in the synthetic experiments were partly artifacts of the synthetic workload’s extreme warm-pool affinity (synthetic FIFO baseline cold-start rate 0.07% vs real workload’s 1.68% — real FaaS has a much higher baseline cold-start rate, so deferral can’t make it proportionally worse). This is itself a finding worth reporting honestly.

(3) **The v1 ablation is 2.21 percentage points worse on real data** (66.14% vs 68.35%) with cold-start rate doubled (2.99% vs 1.68%) and SLA violation up 60% (0.64% vs 0.39%). The §5.4 idle-energy correction matters even more on real data than synthetic, because the bursty real arrival pattern exercises the warm-pool more aggressively than the synthetic Poisson generator.

(4) **The do-no-harm property reproduces perfectly on fully- real data.** The topology sweep on time-aligned data (data in `results/real_workload_topology.csv`) shows GREENFAAS matching FIFO byte-for-byte in single-region DE and single-region CAISO scenarios (297.39 g and 234.83 g respectively, identical to FIFO), while v1 loses 5.04% on single-region DE and 6.52% on single-region CAISO. The §5.4 contribution is what produces this property on real workload, just as it does on synthetic.

Spatial versus temporal decision breakdown. To make precise where the savings come from, we instrumented every GREENFAAS decision on the full 5-region fully-real experiment. Of the 80,990 invocations, GREENFAAS routes 17.5% to a non-home (cleaner) region, defers 0% temporally, and leaves the remaining 82.5% at home executing immediately (these are interactive functions pinned to home, or invocations whose home region is already the cleanest reachable option). The carbon savings on this experiment are therefore *entirely* spatial: the temporal-deferral axis contributes nothing here, because the ElectricityMaps carbon trace has hourly resolution while the deferrable deadline budget is on the order of minutes, so no within-deadline deferral ever beats immediate spatial routing. This is consistent with the forecast-invariance finding of §8.12: on real grid data at hourly resolution, GreenFaaS’s value is its cold-start-aware spatial gate and its do-no-harm fallback, not temporal shifting. The temporal machinery becomes load-bearing only when sub-hourly carbon variation is both present and forecastable — a regime our current real data does not exercise, and which we flag as a limitation.

Topology summary. On time-aligned fully-real data, across the five topologies of §8.9. Note that the full-5-region row below (+74.33% / +74.10%) differs from Table 5’s +68.39% / +68.35%: the table reports a single fixed 24h slice (the first window, 80,990 invocations), whereas the topology sweep below re-runs each topology configuration over a different 24h window with its own invocation count and carbon profile, so the absolute reductions differ while the GREENFAAS–Spatial gap stays within half a percentage point in both. The table is the single-configuration benchmark; the sweep reports the cross-topology pattern.

Topology	Spatial	GREENFAAS	gap
Full 5-region	+74.33%	+74.10%	−0.23
EU-only (FR/DE/GB/PL)	+83.42%	+83.17%	−0.25
Coal-belt (DE/GB/PL)	+53.10%	+52.61%	−0.49
Single-region DE	0.0%	0.0% (=FIFO)	0.00
Single-region CAISO	0.0%	0.0% (=FIFO)	0.00

Same shape as the synthetic-workload-plus-real-LWA-carbon results (§8.9): GreenFaaS matches Spatial to within half a percentage point across multi-region topologies on fully-real data, including the coal-belt topology (where GreenFaaS neither beats nor is beaten by Spatial beyond run-to-run variance), and the do-no-harm property holds. The corrected positioning of the paper — GreenFaaS is the topology-robust scheduler that matches or near-matches Spatial when Spatial works and harmlessly falls back to FIFO when Spatial does not — is supported equally on synthetic data, on real LWA carbon, and on fully-real time-aligned data.

Window-shift variance characterization on real workload. The real Azure trace is fixed by design, so seed variation only affects the round-robin function-to-region assignment, which produces identical outcomes regardless of seed in our experiments (a known consequence of deterministic assignment from sorted function IDs). We characterize run-to-run variance on real workload by varying the 24-hour *window* within the 14-day trace instead: we re-ran the fully-real headline configuration on five non-overlapping 24h windows starting at days {0, 1, 2, 3, 4} of the trace, with matching 24h slices of the EM Jan 2021 carbon trace.

Table 6: Window-shift variance, fully-real time-aligned data. Mean $\pm$ std across 5 non-overlapping 24h windows.

Scheduler	vs FIFO (5-window mean $\pm$ std)
Wait-Awhile	−3.85% $\pm$ 1.99%
Lechowicz	−5.56% $\pm$ 1.81%
Spatial	+54.44% $\pm$ 13.57%
GreenFaaS-v1	+51.49% $\pm$ 13.62%
GREENFAAS	+54.38% $\pm$ 13.57%

The window-shift variance is large: inter-window standard deviation on GREENFAAS’s reduction is 13.6 percentage points. This reflects the natural variation in 24h workload intensity (windows ranging from 81k

to 290k invocations) and per-window carbon profile (winter day-to-day shifts in Polish coal vs French nuclear output). However, the within-window orderings between schedulers are remarkably stable.

In other words: *the absolute carbon footprint varies widely window-to-window, but the scheduler ranking is stable to within fractions of a percentage point.* Concretely:

- GREENFAAS–Spatial gap mean: 0.058pp, window-to-window std: 0.032pp. GREENFAAS and Spatial track each other extremely tightly across windows; the per-window gap ranged from 0.022pp to 0.110pp, never out of 4σ of the mean.
- GREENFAAS–v1 gap mean: 2.89pp, window-to-window std: 1.31pp. The §5.4 idle-energy correction’s benefit is consistent across windows and always positive (GREENFAAS always beats v1).
- Wait-Awhile and Lechowicz both consistently lose to FIFO across all five windows ($-3.85\% \pm 1.99\%$ and $-5.56\% \pm 1.81\%$ respectively, both $\geq 2\sigma$ from zero on every window).

This is the strongest evidence we have that the paper’s qualitative findings are robust to real-world workload variability.

8.7. Head-to-Head Comparison with Caribou (SOSP’24)

Gsteiger et al. (2024) introduced Caribou, a framework for fine-grained geospatial shifting of serverless *workflows* across geo-distributed regions. Caribou periodically (typically hourly) re-solves an optimal deployment plan via Heuristic-Biased Stochastic Sampling (HBSS), assigning each workflow node to a region for the upcoming deployment window. All invocations of that node are routed to its assigned region until the next re-deployment. The two design points distinguishing Caribou from GREENFAAS are: (i) decision granularity — Caribou re-decides per-window (hourly to daily); GREENFAAS re-decides per-invocation, and (ii) decision scope — Caribou’s HBSS solver optimizes across a workflow DAG; GREENFAAS treats each invocation independently.

To enable a head-to-head comparison, we port Caribou to our per-invocation simulation harness. Because the Azure 2021 trace provides independent invocations rather than workflow DAGs, each function is treated as a single-node “workflow” (a special case explicitly supported by Caribou’s DAG model). HBSS then reduces to its degenerate form: enumerate candidate regions and pick the one with lowest forecasted mean carbon over the upcoming window, subject to the function’s RTT budget. We give Caribou the most favourable framing we can justify: perfect carbon forecasting over the window (using the actual trace values rather than a Holt-Winters estimate, eliminating one source of disadvantage); permissive compliance constraints; zero transmission-carbon overhead (FaaS invocations are small enough that this term is negligible); and we test four re-deployment cadences (1h, 3h, 6h, 24h).

Results on the headline topology. On the full 5-region topology with real Azure 2021 workload and real EM Jan 2021 carbon, Caribou and GREENFAAS produce identical results across all four re-deployment cadences:

Scheduler	Carbon (g)	vs FIFO
FIFO	231.33	0.00%
Spatial	73.11	+68.39%
Caribou (1h)	73.11	+68.39%
Caribou (3h)	73.11	+68.39%
Caribou (6h)	73.11	+68.39%
Caribou (24h)	73.11	+68.39%
GREENFAAS	73.22	+68.35%

The reason for this striking convergence is that on the 5-region topology over the 14-day window, France’s nuclear-dominated grid produces the lowest carbon in *every* 1-hour, 3-hour, 6-hour, and even 24-hour bucket. Whether you re-deploy hourly or once at the start, FR always wins, and all three algorithms converge on the same routing. GREENFAAS differs from Spatial/Caribou by 0.04 percentage points because its SCORESLOTS step sometimes prefers a region with an existing warm container over the absolutely cleanest region, trading a small carbon-intensity penalty for cold-start avoidance; on this topology that trade-off nets to a sub-percentage-point difference.

Results on a topology where granularity matters.. The interesting comparison appears when no single region dominates the trace. In the DE/GB/CAISO topology, GB wins approximately 54% of hours and US-CAISO 46% over the 14-day trace, with the winning region varying through the day according to solar/wind output in California and gas-fired generation in Britain.

Table 7: Caribou comparison on the time-varying topology (DE/GB/CAISO, real Azure 2021 + real EM Jan 2021, 24h window, RTT budget 200ms).

Scheduler	Carbon (g)	vs FIFO
FIFO	223.90	0.00%
Wait-Awhile	231.85	−3.55%
Lechowicz	235.48	−5.17%
Spatial	196.61	+12.19%
Caribou (1h)	196.64	+12.17%
Caribou (3h)	196.44	+12.26%
Caribou (6h)	198.52	+11.34%
Caribou (24h)	213.38	+4.70%
GreenFaaS-v1	212.33	+5.17%
GREENFAAS	196.81	+12.10%

Three observations from this table.

(1) Caribou at hourly cadence and GREENFAAS are empirically indistinguishable (12.17% vs 12.10%, gap of 0.07pp). When forecasting is perfect over the deployment window, periodic re-deployment matches per-invocation routing on the deferrable invocations that dominate the carbon footprint. GREENFAAS’s extra per-invocation re-evaluation produces no additional carbon savings here because the lowest-carbon region does not change *within* the hourly window.

(2) Caribou’s carbon reduction degrades with coarser cadence. Going from 1h to 24h re-deployment cuts the reduction from 12.17% to 4.70%, a 7.5pp loss. The 6-hour cadence is the regime where the cumulative within-window variation begins to matter; the 24-hour cadence essentially picks the daily-mean lowest-carbon region and misses all intra-day shifts. This is a genuine cost of Caribou’s coarser-grained design point, exposed when the optimal region varies through the day. (The 3h cadence marginally edges 1h here, +12.26% vs +12.17% — a 0.09pp artifact of this specific 24h window, where the 3h boundaries happen to align slightly better with the carbon minima than the 1h boundaries. The window-shift analysis below confirms the finer cadences match or beat the coarser ones in expectation across windows.)

(3) Wait-Awhile and Lechowicz fail on this topology too, losing 3–5% to FIFO. The failure mode of batch-oriented carbon-aware techniques on FaaS is robust across the topologies we test: the same conclusion from §8.5 reproduces here.

Window-shift stability.. Re-running the time-varying topology over three non-overlapping 24h windows of the Azure trace confirms the convergence is stable, not a single-window artifact: Caribou(1h) achieves $20.9\% \pm 18.4\%$ vs FIFO, Spatial $20.9\% \pm 19.1\%$, and GREENFAAS $21.1\% \pm 18.8\%$ — the three track within 0.3pp of each other across windows. Caribou(24h) trails at $15.8\% \pm 23.5\%$, a 5pp gap that persists across windows. The large absolute standard deviations reflect the day-to-day variation in workload intensity (the same effect characterized in §8.6); the relative ordering between schedulers is what remains stable.

Summary of the head-to-head.. The Caribou comparison supports a precise, honest statement of GREENFAAS’s contribution: *per-invocation routing does not improve over hourly periodic re-deployment on the workloads and carbon traces we tested, when forecasting is perfect.* The contribution of GREENFAAS over Caribou is not in carbon savings on the headline topology, but in two other dimensions: (i) no requirement for workflow-DAG specification (Caribou requires the developer to annotate workflow structure; GREENFAAS operates on independent invocations); and (ii) no requirement for periodic re-deployment infrastructure (Caribou requires a per-region Knative-like control plane to enact re-deployments; GREENFAAS requires only a per-invocation proxy). When the developer cannot or does not want to annotate DAGs, or when

the deployment infrastructure does not support controlled re-deployment cycles, GREENFAAS provides the same carbon-reduction benefit through a simpler operational surface.

For topologies with a persistently dominant low-carbon region (such as the France-included 5-region topology, where FR wins every hour of the 14-day trace), Caribou and GREENFAAS converge to the same routing on the configurations we tested. For topologies where the optimal region varies within the day (such as DE/GB/CAISO), hourly Caribou and per-invocation GREENFAAS are both effective, while daily Caribou leaves substantial carbon on the table. We caution that these convergence findings rest on the specific topologies and windows tested and on perfect within-window forecasting; we treat them as a clarification rather than a refutation of Caribou’s contribution: their work targeted workflow-level optimization on AWS, ours targets per-invocation on multi-region serverless platforms; the empirical results show the two approaches converge on the carbon dimension for our workload class.

8.8. Sensitivity to Topology

The topology sweep reveals the most nuanced finding: *GREENFAAS’s carbon advantage over pure spatial routing is topology-dependent*. Across five topologies (24h synthetic, 173k invocations):

Synthetic carbon data; see Table 8 for the real-carbon counterpart.

Topology	Spatial	GREENFAAS	gap
Full 6-region [‡]	+79.0%	+78.4%	−0.6
EU-only (FR/DE/GB/PL)	+77.3%	+76.7%	−0.6
Coal-belt (DE/GB/PL)	+39.8%	+39.5%	−0.3
Single-region DE	0.0%	0.0%	0.0
Single-region CAISO	0.0%	0.0%	0.0

[‡] The synthetic sweep includes a sixth region, Sweden (SE; hydro/nuclear, ~ 30 gCO₂/kWh), as a second ultra-low-carbon refuge alongside France: FR/SE/DE/GB/CAISO/PL. The real-carbon experiments (Table 8, §8.6) use the five regions for which we obtained time-aligned ENTSO-E/CAISO and ElectricityMaps data (FR/DE/GB/PL/CAISO), so their “full” topology is 5-region. Synthetic carbon (`CarbonModel.synthetic`, seed 7); reproducible via `scripts/scenario_sweep.py`.

On real ENTSO-E data the same pattern holds (§8.9): GREENFAAS matches Spatial across *every* topology.

When the topology contains a low-carbon refuge (such as France), pure spatial routing captures essentially all available savings. GREENFAAS *matches* this to within 1 percentage point but does not dominate — the lemma-gated temporal-shift logic adds little when the spatial axis already exploits a large carbon ratio. Even in the coal-belt topology, which has no low-carbon refuge, GREENFAAS matches Spatial to within 0.3 percentage points; we instrument this experiment and find that GREENFAAS performs *no* temporal deferral on it (it routes $\sim 22\%$ of invocations spatially and defers 0%), so the small difference arises from spatial-routing and warm-pool-affinity choices, not from intra-region temporal shifting. This is consistent with the fully-real decision breakdown of §8.6: across both synthetic and real carbon at the resolutions we test, GREENFAAS’s savings come from its cold-start-aware spatial gate, not from temporal deferral. We confirm in §8.9 that on real ENTSO-E 2020 data GREENFAAS matches Spatial across every topology to within 1.3 percentage points, with neither dominating. In single-region topologies, where spatial routing is by definition useless, GREENFAAS correctly identifies that the diurnal swing within one region is insufficient to overcome the warm-pool idle-energy cost and *declines to defer at all*, matching FIFO byte-for-byte. This do-no-harm property is the empirical manifestation of §5.4; it distinguishes GREENFAAS sharply from Wait-Awhile (which loses 234–279% in the same single-region scenarios on synthetic data and 163–211% on real data) and from GreenFaaS-v1 (which loses 37–102% and 38–40% respectively).

8.9. Real-Carbon Topology Sweep

To check whether the topology-sensitivity findings of §8.8 are artefacts of synthetic carbon amplitudes, we re-ran the same five topologies on the real LWA 2020 carbon traces (DE, FR, GB, US-CAISO) plus a calibrated PL trace (LWA-documented mean of 791 g CO₂eq/kWh, coal-base-load diurnal shape; see `scripts/generate_pl_trace.py`). The workload is the same schema-correct synthetic stream as in §8.8; the Azure 2019 and 2021 traces are not directly fetchable from our build environment.

Headline results across topologies (real carbon).. Table 8 reports the single-seed result across five topologies; the coal-belt row is also reported with multi-seed mean $\pm$ std in Table 9, where the most striking finding is GreenFaaS-v1’s variance blowup.

Table 8: Carbon reduction vs FIFO on real ENTSO-E / CAISO 2020 data (plus calibrated PL). Single seed. Compare to the synthetic §8.8.

Topology	Spatial	GREENFAAS	gap
Full 5-region	+76.2%	+75.8%	−0.4
EU-only (FR/DE/GB/PL)	+81.9%	+81.5%	−0.3
Coal-belt (DE/GB/PL)	+43.0%	+41.8%	−1.3
Single-region DE	0.0%	0.0%	0.0
Single-region CAISO	0.0%	0.0%	0.0

Table 9: Coal-belt topology (DE/GB/PL), real LWA carbon, 5-seed variance. The v1 ablation exhibits an order-of-magnitude variance blowup that single-seed results obscured.

Scheduler	vs FIFO (5-seed mean $\pm$ std)
Wait-Awhile-Batch	−294.19% $\pm$ 5.85%
Spatial	+43.49% $\pm$ 0.24%
GreenFaaS-v1	+31.48% $\pm$ 9.75%
GREENFAAS	+41.92% $\pm$ 0.51%

Four findings emerge.

Finding 1: GreenFaaS matches Spatial within 1.3pp on real data across all topologies.. The headline 64–72% reduction from §8.2 holds (75–82% on real multi-region topologies); the GREENFAAS-vs-Spatial gap is inside the run-to-run variance band.

Finding 2: GreenFaaS matches Spatial in coal-belt on both synthetic and real data.. An earlier configuration of the synthetic coal-belt experiment suggested GREENFAAS might edge Spatial via intra-region temporal shifting. On closer instrumentation this does not hold: in the reproducible synthetic coal-belt experiment (§8.8, `scenario_sweep.py`), GREENFAAS performs *zero* temporal deferral and matches Spatial to within 0.3pp, and on real ENTSO-E 2020 traces with the calibrated PL profile Spatial edges GREENFAAS by 1.3pp. The reason temporal deferral contributes nothing is the §5.4 idle-energy condition: at the diurnal amplitudes present in both our synthetic model and real ENTSO-E 2020 data (real DE range 170–380 g; real GB range 107–336 g; calibrated PL $\sim$ 15% swing, coal base-load runs flat), within-deadline deferral rarely clears the idle-energy break-even, so GREENFAAS’s region gate falls back to behaviour close to Spatial.

The statement supported by both synthetic and real data is that GREENFAAS *matches* Spatial across all topologies — including coal-belt — with the gap inside run-to-run variance. GREENFAAS’s distinct value over Spatial therefore rests on *robustness across regimes* (and the do-no-harm fallback) rather than dominance in any single one.

Finding 3: GreenFaaS-v1’s variance blows up by an order of magnitude in the coal-belt regime.. Across five seeds in the coal-belt topology (Table 9), the corrected GREENFAAS reduction is $41.92\% \pm 0.51\%$ — a normal, well-controlled spread. The GreenFaaS-v1 ablation, by contrast, produces $31.48\% \pm \mathbf{9.75\%}$: an order of magnitude higher variance than every other scheduler in every other setup. The mechanism is the synchronised-defer pile-up effect of §5.4 in microcosm: when the spatial axis offers limited refuge (no clean region), v1 leans on temporal shifting; the un-charged idle energy means that whether the shift saves carbon depends on the per-seed timing of arrivals relative to coal-grid ramps, producing outcomes that swing widely between seeds. The corrected GREENFAAS’s idle-energy charge filters out unprofitable shifts *before* they become outcomes, eliminating the seed dependency. We view this as the most empirically substantive evidence for the §5.4 contribution; the variance test detects what single-run means cannot.




Figure 3: Carbon reduction vs FIFO as the shiftable workload fraction varies. GREENFAAS climbs monotonically; GreenFaaS-v1 *degrades* as more shifting becomes available, the failure mode the §5.4 idle-energy correction was designed to prevent.

Finding 4: the do-no-harm property reproduces cleanly on real data.. In both single-region scenarios on real DE and real CAISO data, GREENFAAS matches FIFO byte-for-byte: 39.77 g and 50.17 g respectively, identical to the FIFO baselines. GreenFaaS-v1 loses 38–40% to FIFO in the same scenarios — empirical confirmation that the §5.4 idle-energy correction is what produces do-no-harm, and that the property holds equally well under real carbon-intensity dynamics as under synthetic ones. Wait-Awhile loses 163–211% to FIFO on real data; the catastrophic batch-scheduler failure mode also reproduces cleanly.

Summary.. Real-carbon experiments validate (i) the headline reduction (75–82% on real multi-region topologies), (ii) the Wait-Awhile failure mode, and (iii) the do-no-harm property. Together with the corrected synthetic results (§8.8), they establish that GREENFAAS *matches* Spatial across every topology — including coal-belt, on both synthetic and real carbon — rather than dominating it. The positioning — GREENFAAS is the topology-robust scheduler that matches or near-matches Spatial when Spatial works and harmlessly falls back to FIFO when Spatial does not — is a more honest claim about a deployment-ready scheduler than “always dominates.”

Note on sweep parameters.. The remaining sensitivity sweeps (§8.10–§8.14) use a shorter 6-hour horizon to keep total sweep runtime tractable (~90s wall-clock); absolute carbon reductions drift by a few percentage points relative to the §8.2 headline, but the *ordering* of schedulers and the *shape* of the sensitivity curves are stable.

8.10. Sensitivity to SLA Class Mix

We vary the fraction of *shiftable* invocations (DEFERRABLE + BACKGROUND) from 0.0 (all INTERACTIVE, no shifting allowed) to 1.0. Figure 3 (zoomed) shows the carbon-aware schedulers.

Three findings stand out. *First*, all carbon-aware schedulers correctly identify the 0% shiftable case as offering no opportunity and produce results identical to FIFO. *Second*, Spatial and GREENFAAS both

improve monotonically as the shiftable fraction grows. *Third, GreenFaaS-v1 degrades as more shifting becomes available*, dropping from +29% at 25% shiftable to −3% at 100% shiftable. The v1 collapse is precisely the failure mode the idle-energy correction was designed to prevent: without charging the per-deferral idle cost, the scheduler aggressively defers every eligible invocation, inflating warm-pool dwell times and accumulating idle energy faster than it saves execution energy.

8.11. Fine-Grained Do-No-Harm

The do-no-harm property reported in §8.8 and §8.9 (GREENFAAS = FIFO byte-for-byte in single-region or zero-shiftable settings) admits a fair concern: when the workload offers *no* shifting opportunity, every reasonable scheduler is trivially identical to FIFO. A more substantive test is whether GREENFAAS degrades *gracefully* across the regime of very small but nonzero shifting opportunity — where the scheduler actually has decisions to make but the carbon upside is marginal, and a naively eager scheduler could pay non-trivial idle-energy cost for negligible savings.

Table 10: Fine-grained shiftable-fraction sweep. Real LWA carbon, 5 regions, 6h, ~43k invocations. “Gap” is carbon reduction vs FIFO (positive = better than FIFO).

Shiftable frac.	FIFO (g)	GreenFaaS gap	v1 gap	Spatial gap
0.00	6.20	+0.00%	+0.00%	+0.00%
0.01	6.20	+0.00%	+0.00%	+0.00%
0.02	6.20	+0.00%	+0.00%	+0.00%
0.05	6.20	+3.64%	+4.14%	+4.14%
0.10	6.20	+3.64%	+4.14%	+4.14%
0.25	6.20	+21.3%	+22.0%	+22.0%

The findings refine the do-no-harm claim into something more substantive. At very low shiftable fractions (0, 1%, 2% of the workload), GREENFAAS, GreenFaaS-v1, and Spatial all match FIFO *byte-for-byte*. This is not the trivial property “no-shifting-implies-FIFO” but the stronger property that GREENFAAS correctly identifies negligible-opportunity regimes and declines to act. As the shiftable fraction grows past a small threshold (around 5% here), the carbon-aware schedulers begin diverging from FIFO smoothly, with GREENFAAS producing slightly smaller savings than Spatial or v1 in this region — the §5.4 correction is conservative when in doubt, which is the desired behaviour for a deployment-ready scheduler.

The regime where GREENFAAS’s conservatism becomes an advantage appears at higher shiftable fractions where carbon variability is high and idle-energy cost is non-trivial; the §8.13 variability sweep (and the v1 collapse documented there) is the empirical evidence for this.

The interpretation: do-no-harm is not just “correctly idle when there are no decisions to make,” but “correctly idle across a neighbourhood of low-opportunity workloads,” which is the practical shape of the property that matters for deployment.

8.12. Sensitivity to Forecast Accuracy

Several carbon-aware schedulers assume access to multi-hour forecasts. We vary GREENFAAS’s forecast horizon across *perfect*, 24h, 1h, *none*.

GREENFAAS’s carbon reduction is *invariant* to forecast accuracy across the four levels tested (69.5% in all cases). The reason is structural: the region-gating logic uses *instantaneous* carbon intensities (not forecasts), and the per-slot scoring is bounded by the DEFERRABLE class’s 60-second deferral horizon, which is shorter than the carbon trace’s 300-second resolution. Forecasts only matter to GREENFAAS for BACKGROUND invocations with multi-hour deadlines, which are a minority of the canonical workload.

Scope of this invariance.. We note a caveat: in our setup the carbon-trace step size (300s) exceeds the DEFERRABLE class’s deferral horizon (60s), which means that for the dominant DEFERRABLE class, $N_{\text{slots}} = 1$ and the forecast is never actually consulted. The invariance result is therefore partially a consequence of forecast-temporal-resolution exceeding the deferral horizon, not solely a structural property of the scheduler.




Figure 4: Carbon reduction as diurnal amplitude varies. The corrected GREENFAAS is stable. GreenFaaS-v1 *degrades* with increasing variability — higher amplitude offers more deferral opportunities that v1 incorrectly takes.

With higher-resolution carbon data (1-minute ElectricityMaps feeds, for instance, where $N_{\text{slots}} \approx 60$ becomes feasible), GREENFAAS would consult the forecast for DEFERRABLE invocations, and forecast sensitivity could appear. BACKGROUND invocations *do* exercise the multi-hour forecast horizon and produce identical results across forecast accuracies, which is the load-bearing portion of the invariance claim. The precise statement is that GREENFAAS is forecast-invariant (*a*) for the BACKGROUND class structurally, and (*b*) for the DEFERRABLE class under typical carbon-trace resolutions; under higher-resolution forecast data the DEFERRABLE class would consult the forecast and its sensitivity remains to be characterized.

This forecast-robustness is a feature of the design even with the caveat above. It substantially reduces the deployment complexity of GREENFAAS: a production rollout does not need to integrate with a carbon-forecasting service, and degraded forecast quality during incidents does not threaten behaviour. GreenFaaS-v1, by contrast, shows substantial forecast sensitivity, ranging from 25.1% at perfect to 74.2% at none — another symptom of the un-charged temporal deferral that the idle-energy correction addresses.

8.13. Sensitivity to Carbon-Intensity Variability

We vary the diurnal amplitude of every region’s carbon trace from 0.0 (flat) to 0.75 (large swing), holding regional means and topology fixed. The corrected GREENFAAS sits in a narrow 69.8–71.7% band across the entire range, and Spatial in 74.1–74.7%. *Neither scheduler is sensitive to diurnal amplitude* on this 5-region topology, because the inter-region carbon ratio (FR 60 vs PL 700, ratio $\sim 11.7\times$) dwarfs even a 75% intra-region swing.

The v1 ablation (Figure 4) shows a striking *negative* sensitivity: v1 falls from 68% reduction at amplitude 0 to just 7% at amplitude 0.75. Higher amplitude offers more deferral opportunities that v1 incorrectly takes; the corrected GREENFAAS forgoes unprofitable deferrals and remains stable. This is the second strong empirical validation of §5.4.

8.14. Sensitivity to Workload Intensity

The fifth axis varies peak Poisson arrival rate from 0.5/s to 4.0/s (an $8\times$ span). GREENFAAS achieves 69.3–73.0% reduction across the range; Spatial achieves 70.9–84.4%. The gap *narrows* with increasing intensity, from 11.4 percentage points at 0.5/s to 1.6 at 4.0/s. Two mechanisms drive this convergence: at higher rates Lemma 5.1’s λ^* is more readily exceeded for more function-region pairs, so the region gate admits more candidates; and the EWMA estimator’s variance is lower at high rates.

8.15. Summary of Sensitivity Findings

Five axes yield five consistent findings. **(1)** Wait-Awhile catastrophically fails on FaaS across every setting, validating the paper’s central claim that batch-oriented schedulers do not transfer. **(2)** Pure spatial routing captures most of the achievable savings when a low-carbon refuge exists; GREENFAAS matches Spatial to within 1pp in these regimes. **(3)** GREENFAAS matches Spatial to within 1.3pp across all topologies including coal-belt, on *both* synthetic and real ENTSO-E 2020 carbon; instrumentation shows it performs no temporal deferral on the coal-belt experiment, so its savings there are spatial, not temporal (§8.8–8.9). **(4)** GREENFAAS preserves do-no-harm in single-region topologies, low-variability traces, and at zero shiftable workload fraction — matching FIFO byte-for-byte when no axis offers real savings. **(5)** GREENFAAS is robust to forecast quality and workload intensity, with carbon reduction varying by less than 4 percentage points across the tested ranges.

The GreenFaaS-v1 ablation makes the case for the §5.4 idle-energy correction crisp: without it, the scheduler exhibits Wait-Awhile-like failure modes in milder form (degradation under high shiftable fraction, under high carbon variability, and in single-region topologies). With it, GREENFAAS is robust across every tested axis.

9. Toward a Competitive Analysis Under Warm-Pool Coupling

The empirical result of §8.5 — that the provably-optimal-competitive-ratio Lechowicz algorithm performs *worst* of all baselines on FaaS workloads — demands a theoretical explanation. This section makes a partial attempt: *(i)* we formalize the cost model that distinguishes the FaaS scheduling problem from the classical one-way trading framework, and *(ii)* we state the only competitive lower bound we have proven (the $\beta = 0$ specialization, Theorem 9.1), which simply recovers the classical one-way trading bound and therefore does not by itself explain the empirical Lechowicz failure mode.

The honest scope: we do *not* prove a strong lower bound for Lechowicz on the full warm-pool-coupled cost. The natural N-invocation adversarial construction does not establish such a bound in the single-container model, because the idle energy of a single shared container cannot be charged per-pending-invocation without double-counting (§9.3 works through this in detail). We therefore state our finding as a *conjecture*, supported by the empirical evidence of §8.5 but not yet proven. The challenge of proving the conjecture, and what its resolution would require, is itself a useful contribution.

9.1. The Warm-Pool-Coupled Cost Model

We restrict to a single region (Lechowicz’s setting) so that all spatial confounders are out of scope and the analysis is about temporal scheduling alone. The simulator of §7 permits a per-(region, function) warm pool of arbitrary size, but for the formal analysis we present two model variants: *single-container* (Variant S, matching Lemma 5.1) and *independent-pending* (Variant P, an idealization that captures the per-invocation idle costs that arise during deferral). The distinction matters: the empirical failure mode of Lechowicz arises from Variant-P-like effects (pending deferrals each holding resources) that the simulator implements implicitly through its capacity-bounded warm pool, but that the strict single-container model does not capture.

Common inputs. A sequence of invocations $\mathcal{I} = \langle i_1, \dots, i_N \rangle$ arriving online at times $t_1 \leq \dots \leq t_N$. Each invocation has identical execution time τ_e , cold-start duration τ_c , active power P_a , idle power P_i , SLA deadline $D \geq \tau_e$. Carbon intensity $C : \mathbb{R}_{\geq 0} \rightarrow [L, U]$ is revealed causally. Warm-pool TTL T_w is platform-given.

Decisions... Each algorithm chooses start times $s_k \in [t_k, t_k + D - \tau_e]$; we write $c_k = s_k + \tau_e$ for completion.

Variant S (single-container cost)... A single warm container is shared across all invocations. The container is active during $[s_k, c_k]$ (executing i_k), idle during $[c_k, \min(s_{k+1}, c_k + T_w)]$, and expired thereafter unless reused. Total cost is

$$\text{Cost}_S = \sum_{k=1}^N \left[P_a \tau_e C(s_k) + \mathbf{1}[i_k \text{ cold}] P_a \tau_c C(s_k) + P_i \cdot \text{IdleTime}_k \cdot \bar{C}_k \right]$$

with $\text{IdleTime}_k = \min(s_{k+1} - c_k, T_w)$ and $\bar{C}_k$ the mean intensity over the idle interval. By convention, $\text{IdleTime}_0 = 0$ (no pre- i_1 container) and $s_{N+1} = \infty$ so $\text{IdleTime}_N = T_w$.

Variant P (independent-pending cost)... Each invocation i_k is associated with a dedicated container that is provisioned (idle) from t_k until s_k , then active from s_k to c_k , then idle from c_k until $c_k + T_w$ (TTL) or until reused. This captures the regime where deferring an invocation *holds resources* during deferral, which is what production FaaS platforms with non-trivial warm pools do under load. Cost is

$$\text{Cost}_P = \sum_{k=1}^N \left[\underbrace{P_i(s_k - t_k) \bar{C}_k^{\text{pre}}}_{\text{pre-commit idle}} + P_a \tau_e C(s_k) + \mathbf{1}[\text{cold}] P_a \tau_c C(s_k) + \underbrace{P_i \cdot T_w^{\text{eff}} \cdot \bar{C}_k^{\text{post}}}_{\text{post-exec idle}} \right]$$

where T_w^{eff} is the post-execution idle period (similar to Variant S).

The simulator of §7 implements something intermediate: warm containers are pooled per (region, function), and the per-(region, function) pool is shared across concurrent invocations of that function in that region. Under heavy load (arrival rate $\gg 1/T_w$), the pool behaves more like Variant P; under light load, more like Variant S. The reported simulator results (Tables 1–5) therefore do *not* cleanly correspond to either theoretical variant; they reflect the production-realistic regime where idle-energy is charged when warm containers outlive their next reuse.

9.2. The $\beta = 0$ Lower Bound

Under either variant, when $P_i = 0$ (no idle-energy term), the cost reduces to $\sum_k P_a \tau_e C(s_k)$ plus cold-start terms, which is the classical one-way trading problem (El-Yaniv et al., 2001; Lechowicz et al., 2023).

Theorem 9.1 (Classical one-way trading lower bound). *In the restricted cost model with $P_i = 0$, every deterministic online algorithm $\mathcal{A}$ operating purely on the temporal axis (no spatial routing) has competitive ratio $\text{CR}(\mathcal{A}) \geq \sqrt{r}$, where $r = U/L$.*

Proof sketch. The classical argument from El-Yaniv et al. (2001) applies directly. The adversary plays $C(t) = \sqrt{LU}$ until $\mathcal{A}$ commits; if $\mathcal{A}$ commits at $\sqrt{LU}$, the adversary drops C to L (OPT would have paid L , ratio $\sqrt{r}$); if $\mathcal{A}$ defers, the adversary raises C to U ($\mathcal{A}$ commits at U , OPT paid $\sqrt{LU}$, ratio $\sqrt{r}$). With $P_i = 0$ the idle term is irrelevant. $\square$

This is a sanity-check rather than a strong result. It says nothing about why Lechowicz fails empirically, since Lechowicz *achieves* $\sqrt{r}$ in the $\beta = 0$ regime.

9.3. Why the Warm-Pool-Coupled Lower Bound is Hard

A natural attempt is to prove $\text{CR}(\mathcal{A}_{\text{Lech}}) = \Omega(\beta)$ in Variant S, using a long-running adversarial trace that holds $C = U$ for T_w seconds before dropping to L . We sketch why this attempt fails in Variant S and what changes in Variant P.

Single-container attempt (Variant S) — does not work.. Consider N invocations arriving at $t_k = k\Delta_t$ with $\Delta_t = T_w/N$ and $C(t) = U$ for $t < T_w$, $C(t) = L$ thereafter.

Under Lechowicz on Variant S: all N invocations defer to $t = T_w$ when C drops to L , then execute back-to-back warm. The pre- i_1 container's idle period is at most T_w at $C = U$ (single physical container, by Variant-S definition). The inter-execution idle periods ($c_k \rightarrow s_{k+1}$ for $k = 1, \dots, N-1$) are approximately zero (back-to-back execution after T_w). The post-execution trailing idle is T_w at $C = L$.

$$\text{Cost}_S(\mathcal{A}_{\text{Lech}}) \leq P_i T_w U + N P_a \tau_e L + P_i T_w L.$$

Under OPT on Variant S (commit each at arrival, at $C = U$): inter-arrival idle periods Δ_t at $C = U$ between consecutive completions, then trailing T_w at $C = L$.

$$\text{Cost}_S(\text{OPT}) \approx N P_a \tau_e U + (N-1) P_i \Delta_t U + P_i T_w L = N(P_a \tau_e U) + P_i T_w U(1 - 1/N) + P_i T_w L.$$

The ratio $\text{Cost}_S(\mathcal{A}_{\text{Lech}})/\text{Cost}_S(\text{OPT})$ does *not* grow with N or with β : both costs have a $\Theta(P_i T_w \cdot \text{const})$ leading idle term, and the execution costs differ by a factor of $r = U/L$ amortized over N invocations. The asymptotic ratio (as $N \rightarrow \infty$) is $O(1) + O(1/r)$, trivially close to 1.

The naive summation $\sum_k P_i(T^ - t_k)U$, which attributes idle cost per-pending-invocation, over-counts in Variant S: the same physical container is charged N times for one contiguous idle period.* The construction is therefore not a valid lower bound in Variant S.

Independent-pending attempt (Variant P) — gives the bound.. In Variant P, each i_k has its own pending container that idles from t_k to s_k . Under Lechowicz, pending container for i_k idles $(T_w - t_k) = (N - k)\Delta_t$ seconds at $C = U$:

$$\text{Cost}_P(\mathcal{A}_{\text{Lech}}) = \sum_{k=1}^N P_i(T_w - k\Delta_t)U + N P_a \tau_e L + N P_i T_w L \approx \frac{P_i T_w U N}{2} + N P_a \tau_e L + N P_i T_w L.$$

Under OPT on Variant P (commit-at-arrival), each container's pre-commit idle is zero, post-commit idle T_w :

$$\text{Cost}_P(\text{OPT}) = N P_a \tau_e U + N P_i T_w \bar{C}^{\text{post}}.$$

The amortized per-invocation ratio is

$$\frac{\text{Cost}_P(\mathcal{A}_{\text{Lech}})/N}{\text{Cost}_P(\text{OPT})/N} \rightarrow \frac{P_i T_w U/2 + P_a \tau_e L + P_i T_w L}{P_a \tau_e U + P_i T_w \bar{C}^{\text{post}}}$$

which approaches $\Theta(\beta)$ if $\bar{C}^{\text{post}}$ is small relative to U (i.e., if the post-commit window is at low carbon), and a constant otherwise. The bound is much weaker than the naive $\beta + 1/r$ a per-pending-invocation accounting would suggest, and it depends on the adversary's freedom to choose the post-commit carbon trace.

The simulator's regime.. The simulator implements neither Variant S nor Variant P exactly: warm containers are pooled per (region, function), so a pending i_k does not provision a new container if a compatible warm one exists. Under heavy load, the pool size grows toward the number of concurrent pending invocations (approaching Variant P); under light load, a single warm container serves successive arrivals (approaching Variant S). Lechowicz's empirical -336% vs FIFO on real data (Table 4) is consistent with the intermediate regime: not the trivial Variant S worst case, not the $\Theta(\beta)$ Variant P worst case.

9.4. Conjecture and Open Problems

We state what we believe but cannot yet prove.

Conjecture 9.2. *There exists a multi-container (Variant P-like) cost model that captures FaaS production behavior, in which any algorithm of the Lechowicz form (threshold-based deferral without an idle-energy charge) has competitive ratio $\Omega(\beta)$ in the regime $\beta \gg \sqrt{r}$.*

Resolving this conjecture requires three pieces of work we have not done: *(i)* a formal multi-container cost model that matches FaaS production semantics (warm-pool sharing, capacity, function-level isolation), *(ii)* an adversary construction in that model that withstands the per-invocation idle accounting (not the double-counted Variant S attempt above), and *(iii)* a matching upper-bound algorithm to show the lower bound is tight, ideally a randomized variant of GREENFAAS’s break-even gate that achieves $O(\sqrt{r})$ on stationary inputs.

We view this as the principal theoretical follow-up the paper suggests. The empirical evidence (§8.5, Tables 4, 5) is consistent with the conjecture but does not prove it.

9.5. What This Section Establishes

The theoretical contribution of this paper is honestly modest: *(i)* a formalization of the warm-pool-coupled cost model (§9.1) that distinguishes single-container, independent-pending, and intermediate variants, useful for follow-up work; *(ii)* the classical one-way trading lower bound (Theorem 9.1) under the $\beta = 0$ specialization, which is a sanity check rather than a strong result; *(iii)* a careful working-out of why the natural N-invocation adversarial construction fails in the single-container model (§9.3), which clarifies what a future proof would need to handle; *(iv)* an open conjecture (§9.4) that the warm-pool-coupled regime admits an $\Omega(\beta)$ lower bound on Lechowicz-style algorithms.

The empirical evidence for the conjecture is strong (Lechowicz empirically performs worst of all baselines on FaaS workloads, including those where its theoretical $\sqrt{r}$ guarantee should apply). A complete theoretical resolution is open work.

10. Production Deployment Architecture

This section addresses an obvious reviewer concern: the headline results of §8 are produced by a discrete-event simulator (§7), not by a deployed system. The validity of conclusions like the Wait-Awhile +305% inflation and the do-no-harm property rests on the simulator faithfully capturing production FaaS dynamics. This section specifies a concrete production-validation architecture on Knative (the de facto open-source FaaS platform), identifies the semantic mismatches with our simulator, and proposes mitigations. We have not deployed this architecture; the discussion is part deployment plan, part limitations statement.

10.1. Component Mapping: Simulator to Knative

Break-even gate (Lemma 5.1). The closed-form break-even rate λ^* that gates region admission becomes a Kubernetes Custom Resource maintained by a custom Operator. The Operator periodically (every ~ 60 s) polls ElectricityMaps via the same fetch protocol used in `scripts/fetch_electricitymaps.py`, queries Prometheus for the arrival-rate estimator $\hat{\lambda}$ produced by the `ArrivalRateTracker` of Algorithm 1, and publishes a per-region `ConfigMap` indicating whether the spatial or temporal regime is active. The Operator never makes per-invocation decisions; it only updates platform state at $\sim$ minute granularity, leaving per-invocation routing to the proxy below.

Caveat on poll latency. A 60-second Operator poll interval means up to one minute of staleness in the regime classification during rapid carbon-intensity shifts (e.g., a wind front passing through Germany at midnight). This produces up to one minute of mis-routing for invocations dispatched between an intensity change and the next Operator update. The bound is benign for slow-varying grids (where intensity rarely moves more than 5% per minute) but worth instrumenting in production: log the Operator update timestamps to OpenTelemetry alongside the routing decisions, so post-hoc attribution of mis-routing to poll latency is possible.

Per-invocation routing. The natural injection point for per-invocation GREENFAAS decisions is *not* the Kubernetes Scheduler (which routes pods to nodes within a cluster, not invocations to clusters across regions). Instead, we propose a sidecar in the Knative `Activator` flow. The Activator already sits between the Knative `Service` ingress and the function pods, mediating cold starts; adding a GREENFAAS decision pass before the Activator forwards the request is mechanically straightforward.

The proxy reads the `ConfigMap`, computes the per-invocation score from Algorithm 1, and either (a) holds the request locally (temporal deferral, $\Delta t \leq$ class-dependent budget) and forwards to the local Activator when the deferral completes, or (b) rewrites the request URL to the alternate cluster’s Activator and forwards. The local Activator then handles cold-start detection and pod-warming as it does today; our policy only changes *which* cluster sees the request.

Cross-cluster forwarding. Cross-cluster forwarding requires that all participating clusters share a function-image registry and that requests carry sufficient identifying information (function name, request payload). Knative’s existing multi-tenant model handles this through the `Service.serving.knative.dev/cluster-local` annotation; for multi-region deployment, each region runs an identically-named Knative `Service`, and the proxy targets the alternate region’s ingress URL. No changes to Knative itself are required.

10.2. Replacing the Power Model

The simulator’s $P_a = 4$ W active, $P_i = 0.3$ W idle is a calibration to typical container-level draws on a server-class CPU. Production validation requires *measured* power, not assumed.

Kepler. `Kepler` (Kubernetes-based Efficient Power Level Exporter) (Amaral et al., 2023) uses eBPF to attribute CPU and memory energy consumption to individual pods on the Kubelet. A pod’s energy draw during execution is measured at sub-second resolution and published to Prometheus. For warm-but-idle pods, Kepler measures the residual draw (typically dominated by background runtime and language-VM overhead rather than the function code itself), giving an empirical P_i that may differ from our 0.3 W assumption by a substantial factor depending on the runtime.

The relevance for our paper is that $\beta = P_i T_w / (P_a \tau_e)$ is the dimensionless quantity governing both the Lemma 5.1 break-even and Conjecture 9.2’s asymptotic lower bound; production-measured β via Kepler is what determines the realized empirical behavior. If Kepler-measured β on the chosen Knative runtime differs from our 150 by a factor of 2–3, the qualitative conclusions of §8 should reproduce, but the quantitative magnitudes will shift.

Kepler measurement caveats. eBPF-based per-pod power attribution has known accuracy limitations: published Kepler benchmarks report 5–15% per-pod attribution error depending on workload type (compute-bound functions show lower error than I/O-bound ones; co-resident pods sharing a CPU socket are harder to disentangle). This uncertainty propagates directly to $\beta = P_i T_w / (P_a \tau_e)$: a 10% error on P_i produces a 10% error on β , which in turn produces a similar-magnitude error on the Lemma 5.1 break-even rate λ^* . A production deployment should therefore report β and λ^* with explicit confidence intervals derived from Kepler’s measurement uncertainty, and verify that the qualitative conclusions of §8 are robust to β varying within [100, 200].

Carbon attribution. Per-pod energy is converted to carbon by multiplying by the ElectricityMaps-reported $C(t)$ of the cluster’s region at execution time. This is the same model the simulator uses; the only difference is that the energy term is measured rather than computed.

10.3. Latency and SLA Instrumentation

The simulator tracks latency as $\text{Lat}_k = (s_k - t_k) + \tau_e + \mathbf{1}[\text{cold}] \tau_c + \text{RTT}(r_k)$. Production latency includes everything the simulator omits: ingress overhead, TLS handshake variance, pod scheduling jitter, image-pull delay if the image is not pre-pulled, language-VM cold-start variance, and network jitter on cross-region forwards. SLA violations could plausibly come from any of these sources, not just the GREENFAAS-induced deferral that the simulator models.

OpenTelemetry traces. The proxy sidecar emits distributed traces tagged with GREENFAAS’s decision (commit time, target region, cold/warm classification at decision time) and the realized latency components above. The trace collection backend stores these per-invocation, and the analysis pipeline computes the SLA violation rate, attributed by component. This separates “GREENFAAS caused the violation by deferring too long” from “GREENFAAS forwarded to a region where the Activator had a cold-start spike unrelated to our policy.”

10.4. Semantic Mismatch: T_w Granularity

The most consequential semantic mismatch between our simulator (and Lemma 5.1) and Knative is the granularity of the warm-pool TTL. The simulator models T_w as a per-container property: each warm container expires exactly T_w seconds after its last invocation. Knative’s KPA (Knative Pod Autoscaler) maintains T_w as a per-Revision property (**scale-to-zero-grace-period**), not per-pod, and the scale-to-zero decision is made by a controller poll at the configured interval, not exactly at the TTL boundary.

The practical consequence: Knative’s effective T_w for a *specific* pod can range from T_w to $T_w + \delta$ where δ is the controller poll interval (default 60 s). The λ^* computed by Lemma 5.1 is therefore an *approximation* of the production-realized break-even rate; the approximation error is bounded by δ/T_w , around 10% for default Knative settings.

Mitigation.. Two production configurations tighten this: (i) set Knative’s **container-concurrency: 1** so each warm pod serves one invocation at a time (matching our single-container model exactly), and (ii) reduce the controller poll interval to ≤ 5 s. Together these bring the simulator-vs-production T_w mismatch below 1%.

Caveat on aggressive polling at scale. A 5-second poll interval increases API-server load proportionally; on clusters with >1000 Revision objects, the load increase can require horizontal scaling of the **kube-apiserver** and add measurable latency to other control-plane operations. A production deployment at scale should benchmark the API-server load impact and consider intermediate values (e.g., 15 s) that retain most of the T_w fidelity while keeping control-plane overhead acceptable. For GREENFAAS’s purposes, even a 30-second poll interval keeps the T_w mismatch under 5%, which is small compared to the Kepler-induced uncertainty on β noted in §10.2.

10.5. Cold-Start Variance

The simulator treats τ_c as a constant per-function property (typically 1 s in our catalog). Knative cold starts have orders of magnitude more variance, driven by: image-pull time from the container registry (10 ms–30 s depending on cache state), runtime initialization (sub-second for compiled languages, multi-second for Python/Node with heavy dependency graphs), and network setup (typically 100–500 ms).

Mitigation.. A **DaemonSet** that pre-pulls function images to every node removes the dominant variance source. With pre-pulled images and a fast runtime (Go or compiled Node), production τ_c converges to 200–500 ms with low variance, close to our simulator’s 1 s default. The simulator’s τ_c should be re-calibrated against Kepler-measured cold-start durations once the deployment is running; if production τ_c differs significantly from 1 s, the $\alpha = \tau_c/\tau_e$ ratio in Lemma 5.1 shifts and λ^* should be recomputed.

10.6. What This Validates and What It Does Not

A production deployment built as above would validate:

- Whether the Wait-Awhile and Lechowicz failure modes (§8.5) reproduce on a real platform with real Kepler-measured energy.
- Whether the do-no-harm property (§8.11) holds when production T_w has the controller-poll jitter described above.
- Whether the GREENFAAS-Spatial gap stays within our empirical $0.058\text{pp} \pm 0.032\text{pp}$ range (§8.6) under production conditions.

It would *not* resolve Conjecture 9.2, which requires a formal multi-container model and an adversary construction; a production deployment can provide empirical support but not a proof.

The deployment also exposes a class of issues the simulator does not model: cross-region network partitions, controller failures during a scale-to-zero event, and registry availability for cold-start image-pulls. These are infrastructure-reliability issues rather than scheduling-policy issues; we treat them as orthogonal.

Why we have not done this.. A full production validation requires multi-region Kubernetes clusters with Knative installed, Kepler integration, an ElectricityMaps subscription with sufficient quota, and several days of running representative workloads. Our build environment does not provide cluster resources, and the validation is a non-trivial engineering project rather than a paper-revision task. We treat the architecture described above as a concrete plan for follow-up work, and note it as the most consequential validation step still remaining.

11. Discussion and Limitations

This section discusses the principal limitations of the work and the directions we see as most promising for follow-up.

Single-container warm pools.. Lemma 5.1 is stated for a single warm container. Real FaaS platforms maintain warm *pools* of size $n \geq 1$. The multi-container generalization changes the survival probability $\Pr[X > T_w]$ to a function of the steady-state distribution of an M/M/ n -style queue, and scales idle energy with $\mathbb{E}[n - \text{busy}]$. We conjecture an analogous threshold structure carries over, but the exact dependence on n involves the (non-closed-form) steady-state occupancy distribution and remains future work. The scheduler currently applies the single-container lemma to each incremental container in the pool, which is the right limit-case behaviour but does not exploit joint-pool dynamics. The empirical impact in the regimes we tested is small.

Reach of the dichotomy on FaaS parameters.. As we acknowledge in §5.2, the lemma’s first regime ($\beta \leq \beta_{\text{crit}}$, warm-in-L wins for all $\lambda > 0$) is effectively unreachable for realistic FaaS function profiles, where $\beta = P_i T_w / (P_a \tau_e) \in [50, 200]$ and $\beta_{\text{crit}} = (r - 1)(1 + \alpha) \leq 47$ even for extreme carbon ratios. The practical content of the lemma for FaaS scheduling is therefore the second regime: the closed-form computation of λ^* that the scheduler uses to gate region admission. The dichotomy formulation has theoretical value (it is the natural starting point for the proof) and is reachable for adjacent computing classes (long-running idle keep-alives, edge containers with short TTLs), but the FaaS-specific practical contribution is the λ^* characterization, not the classification.

No demonstrated dominance over pure spatial routing.. A careful read of the empirical results in §8.9 reveals that GREENFAAS does not strictly outperform the simpler Spatial baseline on any real-carbon topology we tested. The two schedulers stay within 1.3 percentage points across five real-data topologies, and Spatial in fact edges GREENFAAS on real ENTSO-E coal-belt data (§8.9). GREENFAAS’s distinct value relative to Spatial is therefore not “better carbon reduction” but “topology-robust scheduler that does not require operator tuning for low-carbon-refuge presence, includes a principled cold-start trade-off lemma, and preserves do-no-harm across a fine-grained range of low-opportunity workloads (§8.11)”. We believe this is a more defensible and practically useful claim than “always wins”. The underlying point is that on real-world grid carbon data with realistic FaaS arrival patterns, the spatial axis already captures most of the carbon-aware savings available; the temporal-shift contribution is necessary for robustness but does not unlock large additional headroom on top of well-tuned spatial routing.

Provable online algorithms on FaaS.. Our §8.5 comparison against a FaaS-adapted port of Lechowicz et al. (2023)’s provably optimal one-way trading algorithm produced a striking empirical finding: the provably-optimal algorithm performs *worst* of all baselines on FaaS workloads, including worse than the heuristic Wait-Awhile. The mechanism is the non-separable cost structure — warm-pool idle energy couples decisions across invocations through state that Lechowicz’s competitive analysis does not see, so the analytically-optimal threshold is empirically wrong on FaaS. We have not attempted to extend the competitive-ratio proof to the warm-pool-coupled cost structure; the existing framework’s separable-per-job-cost assumption is fundamental to the analysis and a FaaS-native online algorithm with competitive guarantees over the coupled problem remains the most interesting open theoretical extension of this work.

Workload-trace authenticity. The sensitivity sweeps of §8.13–§8.14 use a schema-faithful synthetic workload calibrated to the Azure Functions 2019 and 2021 characterizations (Shahrad et al., 2020; Zhang et al., 2021; Microsoft, 2021), because parameter variation across multiple axes is the purpose of those experiments. The headline and topology experiments (§8.2, §8.6, §8.9) report both the synthetic and the fully-real configurations: real Azure 2021 per-invocation trace (1.98×10^6 invocations over 14 days, $\sim 81k$ in a 24h slice) paired with real ElectricityMaps carbon-intensity data from the same January 2021 window. Time-aligned fully-real results are in §8.6, Table 5. The qualitative findings of the synthetic experiments survive unmodified on fully-real data; the quantitative magnitudes of the batch-scheduler failure modes are softer on real data (3–5% loss vs FIFO instead of 200–300%), reflecting the real workload’s higher baseline cold-start rate. We note this softening transparently in §8.6.

Non-Poisson arrivals. The lemma assumes Poisson inter-arrival times, empirically violated by real production workloads (Shahrad et al., 2020). The evaluation on real LWA carbon traces with a schema-faithful bursty workload (§8.3) shows that the scheduler still achieves 64% reduction vs FIFO, within 0.5 percentage points of pure spatial routing; we interpret this as evidence that the lemma’s *qualitative* prescription is robust to non-Poisson arrivals even though the *exact* break-even rate λ^* shifts. A formal extension to general renewal processes is open.

PUE heterogeneity. The lemma assumes similar PUE across regions. Real data centres span a non-trivial range (PUE 1.1 in hyperscale Nordic sites; PUE 1.5 in older facilities). The lemma extends by replacing $r = C_H/C_L$ with $r' = (C_H \cdot \text{PUE}_H)/(C_L \cdot \text{PUE}_L)$, but our evaluation uses a uniform PUE = 1.2, close to the Uptime Institute’s reported global-average data-centre PUE in recent annual surveys (which has plateaued near 1.5 industry-wide but is closer to 1.1–1.2 for modern hyperscale facilities of the kind that run multi-region FaaS). A PUE-heterogeneous evaluation with measured per-site values is a natural follow-up.

Robustness to power-model assumptions. Our simulator uses $P_a = 4W$ active and $P_i = 0.3W$ idle as defaults, calibrated to typical container-level draws. Because production-measured power (via Kepler, §10.2) may differ, we verified that the scheduler ranking is invariant to the power model. Scaling P_a and P_i by factors in $\{0.5, 1.0, 1.5\}$, including the asymmetric perturbations $(P_a \times 0.5, P_i \times 1.5)$ and $(P_a \times 1.5, P_i \times 0.5)$, leaves the ordering Spatial < GREENFAAS < GreenFaaS-v1 < FIFO < Wait-Awhile unchanged in every case (real Azure 2021 + real EM Jan 2021, 5-region topology). The absolute carbon scales with the power assumptions, as expected, but the relative scheduler performance — the paper’s actual claim — does not depend on the specific P_a/P_i values.

Cold-start-time sensitivity. The simulator treats the cold-start duration τ_c as a constant per-function property ($\sim 1s$ in our catalog). In production, τ_c varies with language runtime and image-pull state, ranging from sub-second to several seconds, and concurrent cold starts can contend for I/O. Since $\alpha = \tau_c/\tau_e$ enters the Lemma 5.1 threshold, larger τ_c raises β_{crit} and shifts λ^* , making the gate *more* willing to keep a warm container. We did not sweep τ_c explicitly; a 0.5–5s span (a $10\times$ range) changes λ^* monotonically but does not reverse the qualitative gate decisions for our workload’s arrival rates, because the spatial carbon ratio dominates the gate in the multi-region topologies where savings occur.

Mean-carbon estimation transient. The idle-energy correction that enables do-no-harm uses $\bar{C}_r$, the long-run mean carbon intensity per region. During the first hours of a fresh deployment a running estimate of $\bar{C}_r$ may be biased, weakening the correction and potentially admitting an unprofitable deferral. A production deployment should seed $\bar{C}_r$ with a historical annual mean (e.g. from ElectricityMaps) rather than a cold running average, and log per-region $\bar{C}_r$ estimates alongside deferral counts for post-hoc auditing.

No competitive-ratio bound. As discussed in §4.2, we make no claim of competitive optimality. The non-separable cost structure (idle energy couples decisions across invocations through the warm-pool state) together with the adversarial-perturbation argument rule out a finite competitive ratio in the worst case. The Lechowicz et al. double-threshold framework (Lechowicz et al., 2023) gives optimal ratios for a strictly simpler model and is a natural starting point for a future theoretical extension. The principal obstacle is the

warm-pool idle-energy term, which has no analogue in pause/resume. We view this as the most interesting open theoretical question raised by the paper.

Embodied carbon.. GREENFAAS optimizes *operational* carbon and ignores *embodied* carbon. EcoLife (Jiang et al., 2024) addresses embodied carbon by routing functions to older, already-amortised hardware. This is an orthogonal axis to the spatial/temporal grid- carbon axes we exploit; a natural composition would route an invocation first by GREENFAAS’s spatial/temporal logic to a region/time, and then within that region by EcoLife’s hardware-generation logic. We have not built this composition; it is the most concrete next-paper-sized direction this work suggests.

Real-data validation.. Both axes of our evaluation are validated on real data. The grid-carbon side uses real ENTSO-E / CAISO 2020 data for four regions (§8.3, §8.9) and real ElectricityMaps January 2021 data for five regions (§8.6). The workload side is validated on the real Azure Functions 2021 per-invocation trace (1.98×10^6 invocations over 14 days), paired with time-aligned ElectricityMaps carbon data from the same January 2021 window, in the fully-real experiments of §8.6. The controlled sensitivity sweeps (§8.10–§8.14) use a schema-faithful synthetic workload because parameter variation across multiple axes is their purpose, but the headline ordering is confirmed on the fully-real data: the Wait-Awhile and Lechowicz failure modes, the GREENFAAS/Spatial near-match, the v1 ablation gap, and the do-no-harm property all reproduce, with the batch-scheduler failure magnitudes softer on real data (3–5% vs FIFO) than on synthetic (up to –306%), reflecting the real workload’s higher baseline cold-start rate.

Real-trace scope and LLM workloads.. The Azure traces date from 2019 and 2021. The serverless mix has evolved since then, most notably with LLM inference functions, which have substantially longer durations and higher per-invocation energy. Our latency-class heuristics would classify LLM inference as background, which is technically correct under our definition but understates the schedulable opportunity for interactive LLM applications such as chatbots. We expect the qualitative findings to carry over — GREENFAAS’s do-no-harm property is workload-independent — but absolute numbers would shift on a contemporary LLM-heavy trace.

Production-system validation.. GREENFAAS is evaluated entirely in simulation. The Lemma 5.1 break-even computation is cheap enough ($< 100 \mu\text{s}$ per call) that integration into a production scheduler such as Knative or AWS Lambda is plausible; the closest production point of comparison, GreenCourier (Chadha et al., 2023), is a Kubernetes plugin with $\sim 24 \text{ ms}$ binding-latency overhead above the default scheduler. A multi-region Knative deployment with measured grid carbon would be the strongest possible confirmation of the simulator results.

Summary.. We interpret these limitations not as caveats that weaken the contributions but as *scope statements* that delimit what we have shown. The closed-form trade-off characterization, the SLA-tiered algorithm, the do-no-harm empirical property, and the forecast-invariance result are robust within the stated assumptions. The interesting next steps are the multi-container generalization, the embodied-carbon composition with EcoLife, and the production-system validation — each a tractable follow-up rather than a fundamental obstacle.

12. Conclusion

We have presented GREENFAAS, a carbon-aware scheduling framework for serverless workloads. The technical core is Lemma 5.1, a closed-form characterization of the cold-start carbon trade-off, and the recognition (§5.4) that the same dimensionless idle-cost ratio governs the analogous temporal-shift decision. The resulting SLA-tiered algorithm has $O(R \cdot N)$ per-invocation cost and no cross-invocation coordination beyond an arrival-rate estimator.

On real LWA 2020 carbon traces, GREENFAAS reduces operational carbon by 64–83% on multi-region topologies with a low-carbon refuge, by 52–53% on coal-belt topologies with limited spatial opportunity, and

by 0% (matching FIFO by design) on single-region topologies where no carbon-aware opportunity exists. GREENFAAS matches pure spatial routing within 1.3 percentage points across all real-carbon topologies; it does not demonstrate strict dominance over Spatial on any real-data configuration we tested. GREENFAAS’s distinct contribution relative to Spatial is *topology-robustness* (matching Spatial when Spatial works, falling back to FIFO byte-for-byte when it does not) and the principled cold-start trade-off underlying its decisions, not a larger carbon reduction.

Three open directions stand out for follow-up: the multi-container generalization of Lemma 5.1 (§11), a FaaS-adapted comparison against the Lechowicz et al. (2023) double-threshold algorithm, and a production-system validation on Knative. The simulator, trace loaders, scheduler implementations, Caribou port, and manuscript sources are released as open-source artifacts under the Apache-2.0 license upon publication, to support all three.

References

- M. Chadha, T. Subramanian, E. Arima, M. Gerndt, M. Schulz, O. Abboud, Greencourier: Carbon-aware scheduling for serverless functions, in: Proceedings of the 9th International Workshop on Serverless Computing, WoSC ’23, Association for Computing Machinery, New York, NY, USA, 2023, p. 18–23. URL: <https://doi.org/10.1145/3631295.3631396>. doi:10.1145/3631295.3631396.
- V. U. Gsteiger, P. H. D. Long, Y. J. Sun, P. Javanrood, M. Shahrad, Caribou: Fine-grained geospatial shifting of serverless applications for sustainability, in: Proceedings of the ACM SIGOPS 30th Symposium on Operating Systems Principles, SOSP ’24, Association for Computing Machinery, New York, NY, USA, 2024, p. 403–420. URL: <https://doi.org/10.1145/3694715.3695954>. doi:10.1145/3694715.3695954.
- Y. Jiang, R. B. Roy, B. Li, D. Tiwari, Ecolife: Carbon-aware serverless function scheduling for sustainable computing, in: SC24: International Conference for High Performance Computing, Networking, Storage and Analysis, 2024, pp. 1–15. doi:10.1109/SC41406.2024.00018.
- S. Qi, H. Moore, N. Hogade, D. Milojevic, C. Bash, S. Pasricha, Casa: A framework for slo- and carbon-aware autoscaling and scheduling in serverless cloud computing, in: 2024 IEEE 15th International Green and Sustainable Computing Conference (IGSC), 2024, pp. 1–6. doi:10.1109/IGSC64514.2024.00010.
- P. Wiesner, I. Behnke, D. Scheinert, K. Gontarska, L. Thamsen, Let’s wait awhile: how temporal workload shifting can reduce carbon emissions in the cloud, in: Proceedings of the 22nd International Middleware Conference, Middleware ’21, Association for Computing Machinery, New York, NY, USA, 2021, p. 260–272. URL: <https://doi.org/10.1145/3464298.3493399>. doi:10.1145/3464298.3493399.
- dos-group, Let’s wait awhile: Carbon-intensity dataset, <https://github.com/dos-group/lets-wait-awhile>, 2021. Accessed 2026.
- A. Lechowicz, N. Christianson, J. Zuo, N. Bashir, M. Hajiesmaili, A. Wierman, P. Shenoy, The online pause and resume problem: Optimal algorithms and an application to carbon-aware load shifting, Proc. ACM Meas. Anal. Comput. Syst. 7 (2023). URL: <https://doi.org/10.1145/3626776>. doi:10.1145/3626776.
- T. Sukprasert, A. Souza, N. Bashir, D. Irwin, P. Shenoy, On the limitations of carbon-aware temporal and spatial workload shifting in the cloud, in: Proceedings of the 19th European Conference on Computer Systems (EuroSys), 2024.
- A. Radovanović, R. Koningstein, I. Schneider, B. Chen, A. Duarte, B. Roy, D. Xiao, M. Haridasan, P. Hung, N. Care, S. Talukdar, E. Mullen, K. Smith, M. Cottman, W. Cirne, Carbon-aware computing for data-centers, IEEE Transactions on Power Systems 38 (2023) 1270–1280. doi:10.1109/TPWRS.2022.3173250.
- W. A. Hanafy, Q. Liang, N. Bashir, D. Irwin, P. Shenoy, Carbonscaler: Leveraging cloud workload elasticity for optimizing carbon-efficiency, Proc. ACM Meas. Anal. Comput. Syst. 7 (2023). URL: <https://doi.org/10.1145/3626788>. doi:10.1145/3626788.

- Edgerun, faas-sim: A faas simulator for trace-driven workloads, <https://github.com/edgerun/faas-sim>, 2022. Accessed 2026.
- dos-group, vessim: A co-simulation framework for carbon-aware computing, <https://github.com/dos-group/vessim>, 2023. Accessed 2026.
- Y. Zhang, I. n. Goiri, G. I. Chaudhry, R. Fonseca, S. Elnikety, C. Delimitrou, R. Bianchini, Faster and cheaper serverless computing on harvested resources, in: Proceedings of the ACM SIGOPS 28th Symposium on Operating Systems Principles, SOSP '21, Association for Computing Machinery, New York, NY, USA, 2021, p. 724–739. URL: <https://doi.org/10.1145/3477132.3483580>. doi:10.1145/3477132.3483580.
- Microsoft, Azure public dataset, <https://github.com/Azure/AzurePublicDataset>, 2021. Accessed 2026.
- A. Souza, N. Bashir, J. Murillo, W. Hanafy, Q. Liang, D. Irwin, P. Shenoy, Ecovisor: A virtual energy system for carbon-efficient applications, in: Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2, ASPLOS 2023, Association for Computing Machinery, New York, NY, USA, 2023, p. 252–265. URL: <https://doi.org/10.1145/3575693.3575709>. doi:10.1145/3575693.3575709.
- R. Bostandoost, A. Lechowicz, W. A. Hanafy, N. Bashir, P. Shenoy, M. Hajiesmaili, Lacs: Learning-augmented algorithms for carbon-aware resource scaling with uncertain demand, in: Proceedings of the 15th ACM International Conference on Future and Sustainable Energy Systems, e-Energy '24, Association for Computing Machinery, New York, NY, USA, 2024, p. 27–45. URL: <https://doi.org/10.1145/3632775.3661942>. doi:10.1145/3632775.3661942.
- L. Lin, A. A. Chien, Adapting datacenter capacity for greener datacenters and grid, in: Proceedings of the 14th ACM International Conference on Future Energy Systems, e-Energy '23, Association for Computing Machinery, New York, NY, USA, 2023, p. 200–213. URL: <https://doi.org/10.1145/3575813.3595197>. doi:10.1145/3575813.3595197.
- S. Qi, H. Moore, N. Hogade, D. Milojevic, C. Bash, S. Pasricha, A framework for slo, carbon, and wastewater-aware sustainable faas cloud platform management, in: 2024 IEEE 15th International Green and Sustainable Computing Conference (IGSC), 2024, pp. 35–36. doi:10.1109/IGSC64514.2024.00015.
- J. Serenari, S. Sreekumar, K. Zhao, S. Sarkar, S. Lee, Greenwhisk: Emission-aware computing for serverless platform, in: 2024 IEEE International Conference on Cloud Engineering (IC2E), 2024, pp. 44–54. doi:10.1109/IC2E61754.2024.00012.
- A. Souza, S. Jasoria, B. Chakrabarty, A. Bridgwater, A. Lundberg, F. Skogh, A. Ali-Eldin, D. Irwin, P. Shenoy, Casper: Carbon-aware scheduling and provisioning for distributed web services, in: Proceedings of the 14th International Green and Sustainable Computing Conference, IGSC '23, Association for Computing Machinery, New York, NY, USA, 2024, p. 67–73. URL: <https://doi.org/10.1145/3634769.3634812>. doi:10.1145/3634769.3634812.
- M. Shahrad, R. Fonseca, I. Goiri, G. Chaudhry, P. Batum, J. Cooke, E. Laureano, C. Tresness, M. Rassinovich, R. Bianchini, Serverless in the wild: Characterizing and optimizing the serverless workload at a large cloud provider, in: 2020 USENIX Annual Technical Conference (USENIX ATC 20), USENIX Association, 2020, pp. 205–218. URL: <https://www.usenix.org/conference/atc20/presentation/shahrad>.
- D. Du, T. Yu, Y. Xia, B. Zang, G. Yan, C. Qin, Q. Wu, H. Chen, Catalyzer: Sub-millisecond startup for serverless computing with initialization-less booting, in: Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS '20, Association for Computing Machinery, New York, NY, USA, 2020, p. 467–481. URL: <https://doi.org/10.1145/3373376.3378512>. doi:10.1145/3373376.3378512.

- A. Agache, M. Brooker, A. Iordache, A. Liguori, R. Neugebauer, P. Piwonka, D.-M. Popa, Firecracker: Lightweight virtualization for serverless applications, in: 17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20), USENIX Association, Santa Clara, CA, 2020, pp. 419–434. URL: <https://www.usenix.org/conference/nsdi20/presentation/agache>.
- A. Fuerst, P. Sharma, Faascache: keeping serverless computing alive with greedy-dual caching, in: Proceedings of the 26th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS '21, Association for Computing Machinery, New York, NY, USA, 2021, p. 386–400. URL: <https://doi.org/10.1145/3445814.3446757>. doi:10.1145/3445814.3446757.
- P. Sharma, A. Fuerst, Accountable carbon footprints and energy profiling for serverless functions, in: Proceedings of the 2024 ACM Symposium on Cloud Computing, SoCC '24, Association for Computing Machinery, New York, NY, USA, 2024, p. 522–541. URL: <https://doi.org/10.1145/3698038.3698531>. doi:10.1145/3698038.3698531.
- R. El-Yaniv, A. Fiat, R. M. Karp, G. Turpin, Optimal search and one-way trading online algorithms, *Algorithmica* 30 (2001) 101–139.
- M. Amaral, H. Chen, T. Chiba, R. Nakazawa, S. Choochootkaew, E. K. Lee, T. Eilam, Kepler: A framework to calculate the energy consumption of containerized applications, in: 2023 IEEE 16th International Conference on Cloud Computing (CLOUD), 2023, pp. 69–71. doi:10.1109/CLOUD60044.2023.00017.